


# Reinforcement Learning

## Chapter 21



**Mausam**

(some slides by Rajesh Rao)



**1**





**2**



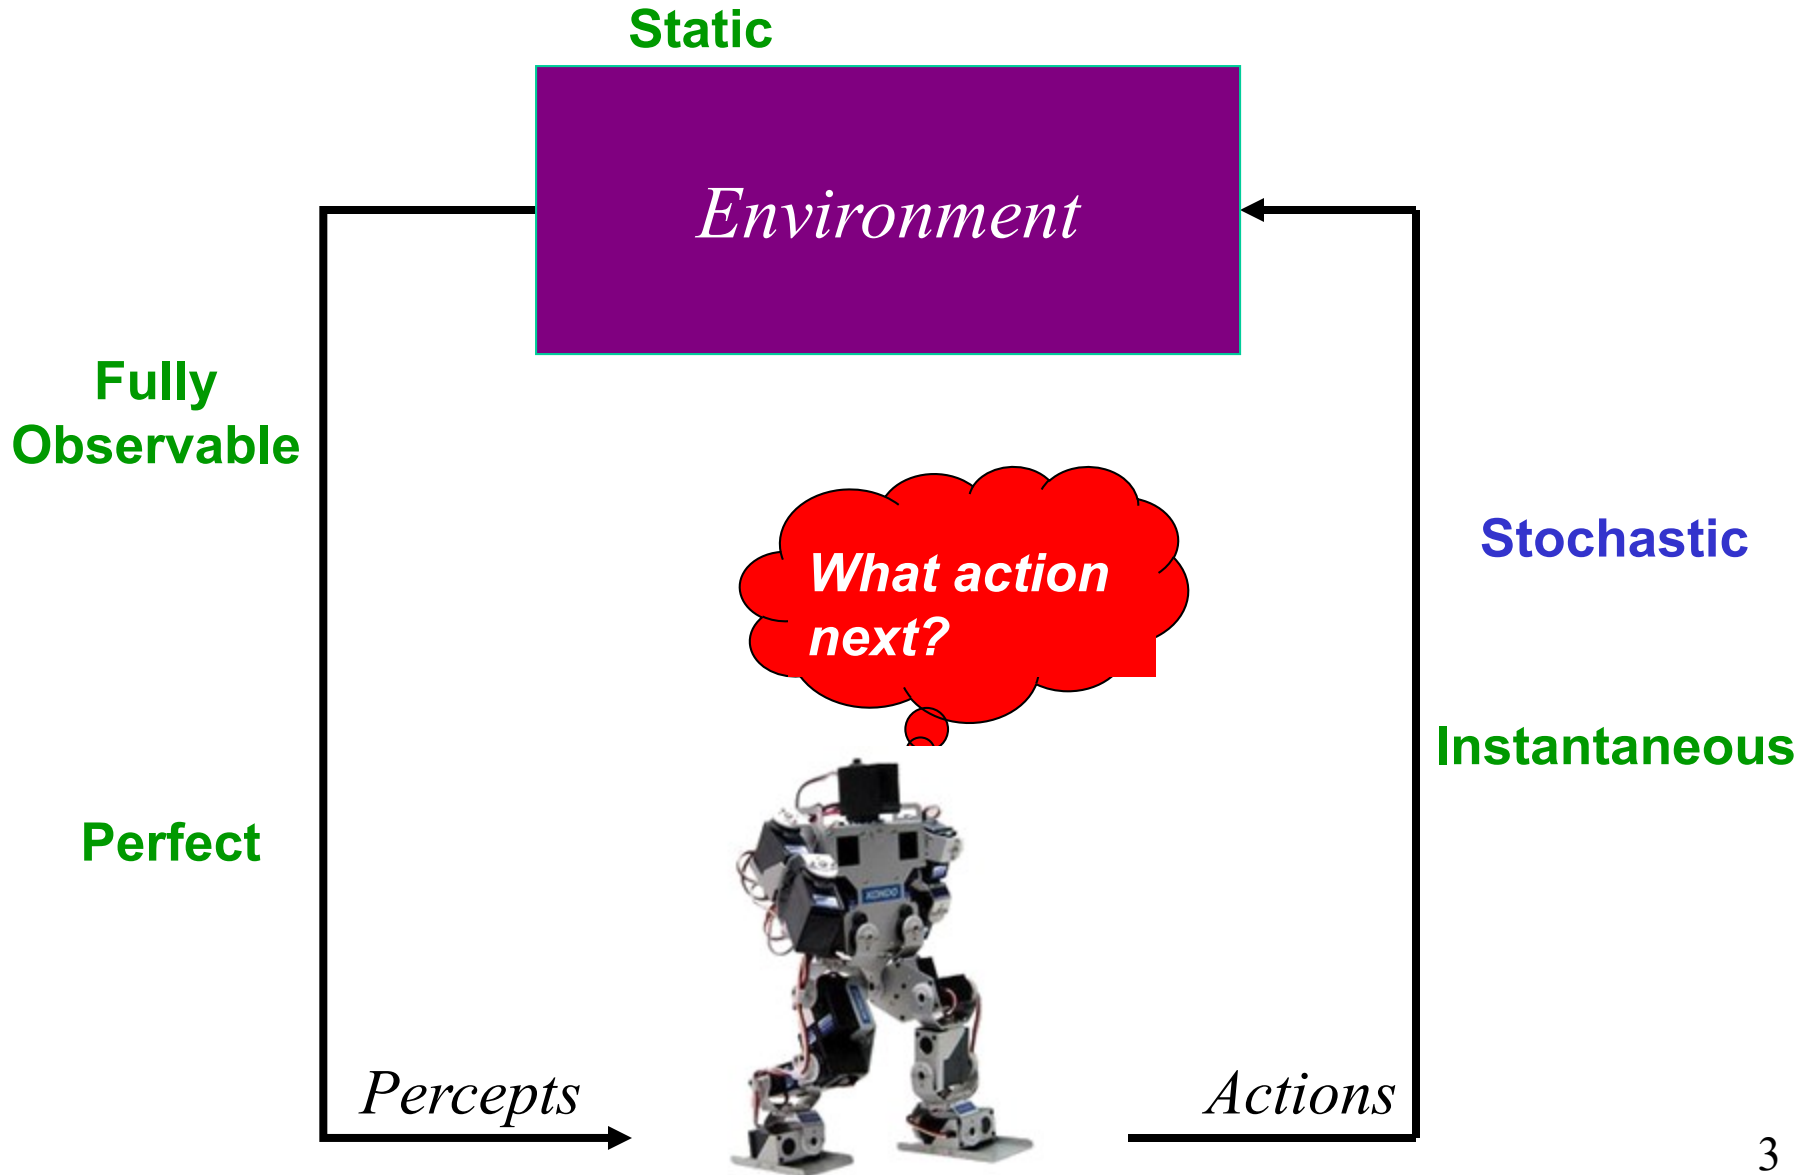
**3**



**4**



# MDPs



# Reinforcement Learning

- $S$ : a set of states
- $A$ : a set of actions
- $T(s,a,s')$ : transition model
- $R(s,a)$ : reward model
- $\gamma$ : discount factor
- Still looking for policy  $\pi(s)$

- New Twist: we don't know  $T$  and/or  $R$ 
  - we don't know which state is good/what actions do
  - must learn from data/experience
- Fundamental model for learning of human behavior

# Learning vs Inference

- Batch setting in Bayes Nets
  - Data  $\rightarrow$  Model  $\rightarrow$  Prediction


- Active setting in MDPs
  - Action  $\rightarrow$  Data  $\rightarrow$  (Model?)



- Actions have two purposes
  - To maximize reward
  - To learn the model

# Learning/Planning/Acting

---



# Main Dimensions

## ■ Model-based vs. Model-free

- Model-based: learn the model ( $T$ ,  $R$ )
- Model-free: directly learn what action to do when

## ■ Passive vs. Active

- Passive: learn state values evaluating a given policy
- Active: need to learn both optimal policy + state values

## ■ Strong vs Weak simulator

- Strong: can jump to any part of state space and simulate
- Weak: real world; can't teleport

# RL and Animal Foraging

- RL studied experimentally for more than 80 years in psychology and brain science
  - Rewards: food, pain, hunger, drugs, etc.
  - Evidence for RL in the brain via a chemical called dopamine
- Example: foraging
  - Bees can learn near-optimal foraging policy in field of artificial flowers with controlled nectar supplies



# Passive Learning (Policy Evaluation)

- Given a policy  $\pi$ : compute  $V^\pi$ 
  - $V^\pi$  : expected discounted reward while following  $\pi$
- Remember
  - We don't know  $T$
  - We don't know  $R$
  - But we can execute (and simulate)
- Key Idea
  - compute expectations by average over samples

## Aside: Expected Age

Goal: Compute expected age of COL333 students

Known  $P(A)$

$$E[A] = \sum_a P(a) \cdot a$$

$$35 \times 20 + \dots$$

Without  $P(A)$ , instead collect samples  $[a_1, a_2, \dots, a_N]$

Unknown  $P(A)$ : “Model Based”

$$\hat{P}(a) = \frac{\text{num}(a)}{N}$$

$$\approx \sum_a \hat{P}(a) \cdot a$$

Why does this work? Because eventually you learn the right model.

Unknown  $P(A)$ : “Model Free”

$$\approx \frac{1}{N} \sum_i a_i$$

Why does this work? Because samples appear with the right frequencies.

## Method 1: Model-based Learning

- Learn an empirical model
- Solve for  $V^\pi$  using policy evaluation
  - assuming that the learned model is correct
- Learning the model
  - maintain estimates of  $T(s,a,s')$
  - maintain estimates of  $R(s,a,s')$

## Example

- 12 states, 4 actions
- $\text{Reward}(\text{action}) = -1$
- Discount factor = 1
- A4 and C4 are absorbing states
- When might this be the optimal policy?

	1	2	3	4
A	↓	↓	→	+100
B	→	→	↑	←
C	→	→	↑	-100

## Data on Executing $\pi$

(A1, D, -1) (A1, D, -1)  
(B1, R, -1) (B1, R, -1)  
(B2, R, -1) (B2, R, -1)  
(B3, U, -1) (B3, U, -1)  
(A3, R, -1) (C3, U, -1)  
(A2, D, -1) (C4, -100)  
(B2, R, -1)  
(B3, U, -1)  
(A3, R, -1)  
(A4, 100)

	1	2	3	4
A	↓	↓	→	+100
B	→	→	↑	←
C	→	→	↑	-100

- $T(A1, D, B1) = 1$
- $T(B3, U, A3) = 2/3$
- We may want to smooth...

# Properties

- Converges to correct model with infinite data
  - If no state is starved
- With correct model
  - $V^\pi$  is computed accurately
- How about model free learning?
  - i.e., expectation is average of samples

## Method 2: Empirical Estimation of $V^\pi$

- Given a policy  $\pi$ : compute  $V^\pi$ 
  - $V^\pi$  : expected discounted long-term reward following  $\pi$
  - $V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [\textit{long term reward with } s \rightarrow s']$
  - $V^\pi(s) = \frac{1}{N} \sum_i [\textit{long term reward}_i]$

## Data on Executing $\pi$

(A1, D, -1)    (A1, D, -1)  
 (B1, R, -1)    (B1, R, -1)  
 (B2, R, -1)    (B2, R, -1)  
 (B3, U, -1)    (B3, U, -1)  
 (A3, R, -1)    (C3, U, -1)  
 (A2, D, -1)    (C4, -100)  
 (B2, R, -1)  
 (B3, U, -1)  
 (A3, R, -1)  
 (A4, 100)

	1	2	3	4
A	↓	↓	→	+100
B	→	→	↑	←
C	→	→	↑	-100

- $V^\pi(B1) =$
- $V^\pi(B2) =$



# Properties

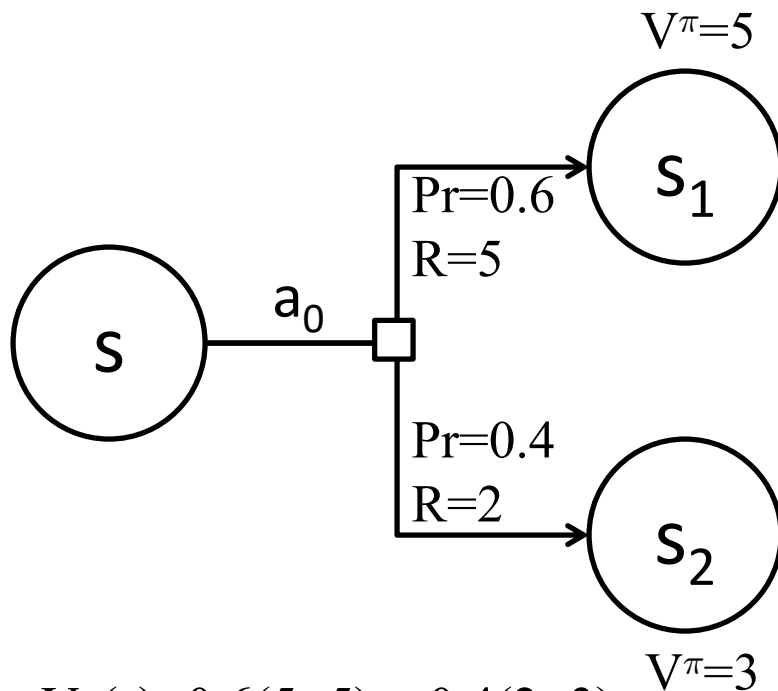
- Converges to optimal with infinite data
  - If no state is starved
- Is wasteful (why?)
  - Compare  $V^\pi(B1)$  and  $V^\pi(B2)$
- Each state is computed independently
  - Connections (Bellman equations) are ignored
  - Learns slowly

## Method 3: Temporal Difference Learning


- Given a policy  $\pi$ : compute  $V^\pi$ 
  - $V^\pi$  : expected discounted long-term reward following  $\pi$
  - $V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [\textit{long term reward with } s \rightarrow s']$
  - $V^\pi(s) = \frac{1}{N} \sum_i [\textit{long term reward}_i]$
- $V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$
- represents relationship between  $s$  and  $s'$
- TD Learning: computing this expectation as average

# TD Learning

- $V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$
- Say I know correct values of  $V^\pi(s_1)$  and  $V^\pi(s_2)$



$$\begin{aligned} V^\pi(s) &= 0.6(5+5) + 0.4(2+3) \\ &= 6 + 2 = 8 \end{aligned}$$



$$\begin{aligned} V^\pi(s) &= (10+10+10+5+5)/5 \\ &= 8 \end{aligned}$$

# TD Learning

- $V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$
- Inner term is the sample value
  - (s,s',r): reached s' from s by executing  $\pi(s)$  and got immediate reward of r
  - sample =  $r + \gamma V^\pi(s')$
- Compute  $V^\pi(s) = \frac{1}{N} \sum_i sample_i$
- Problem: we don't know true values of  $V^\pi(s')$ 
  - learn together using dynamic programming!

## Estimating mean via online updates

- Don't learn T or R; directly maintain  $V^\pi$
- Update  $V^\pi(s)$  each time you take an action in s via a moving average

- $V_{n+1}^\pi(s) \leftarrow \frac{1}{n+1} (n \cdot V_n^\pi(s) + \text{sample}_{n+1})$

- $V_{n+1}^\pi(s) \leftarrow \frac{1}{n+1} ((n+1-1) \cdot V_n^\pi(s) + \text{sample}_{n+1})$

- $V_{n+1}^\pi(s) \leftarrow V_n^\pi(s) + \frac{1}{n+1} (\text{sample}_{n+1} - V_n^\pi(s))$

average of n+1 samples

learning rate

sample n+1

- $V_{n+1}^\pi(s) \leftarrow V_n^\pi(s) + \alpha (\text{sample}_{n+1} - V_n^\pi(s))$

- Nudge the old estimate towards the sample

# TD Learning

- $(s, s', r)$
- $V^\pi(s) \leftarrow V^\pi(s) + \alpha(\text{sample} - V^\pi(s))$
- $V^\pi(s) \leftarrow V^\pi(s) + \alpha(r + \gamma V^\pi(s') - V^\pi(s))$  TD-error
- $V^\pi(s) \leftarrow (1 - \alpha)V^\pi(s) + \alpha(r + \gamma V^\pi(s'))$
- Update maintains a mean of (noisy) value samples
- If the learning rate decreases appropriately with the number of samples (e.g.  $1/n$ ) then the value estimates will converge to true values! (non-trivial)

# Early Results: Pavlov and his Dog

- Classical (Pavlovian) conditioning experiments
- Training: Bell → Food
- After: Bell → Salivate
- Conditioned stimulus (bell) predicts future reward (food)




FIG. 2.

## Predicting Delayed Rewards



- Reward is typically delivered at the end (when you know whether you succeeded or not)
- Time:  $0 \leq t \leq T$  with stimulus  $a(t)$  and reward  $r(t)$  at each time step  $t$  (Note:  $r(t)$  can be zero at some time points)
- Key Idea: Make the output  $v(t)$  predict total expected future reward starting from time  $t$

$$v(t) \approx \left\langle \sum_{\tau=0}^{T-t} r(t + \tau) \right\rangle$$



# Predicting Delayed Reward: TD Learning


Stimulus at  $t = 100$  and reward at  $t = 200$



Prediction error  $\delta$  for each time step  
(over many trials)


# Prediction Error in the Primate Brain?

## Dopaminergic cells in Ventral Tegmental Area (VTA)



# More Evidence for Prediction Error Signals

## Dopaminergic cells in VTA



Negative error

$$r(t) = 0, v(t+1) = 0$$

$$[r(t) + v(t+1) - v(t)] = -v(t)$$

# The Story So Far: MDPs and RL

## Known MDP: Offline Solution

### Goal

Compute  $V^*, Q^*, \pi^*$   
Evaluate a policy  $\pi$

### Technique

Value / policy iteration  
Policy evaluation

## Unknown MDP: Model-Based

### Goal

Compute  $V^*, Q^*, \pi^*$   
Evaluate a policy  $\pi$

### Technique

VI/PI on approx. MDP  
PE on approx. MDP

## Unknown MDP: Model-Free

### Goal

Compute  $V^*, Q^*, \pi^*$   
Evaluate a policy  $\pi$

### Technique

Q-learning  
TD-Learning

# Model-based RL

- Learn an initial model  $M_0$
- Loop
  - VI/PI on  $M_i$  to compute policy  $\pi_i$
  - Execute  $\pi_i$  to generate data
  - Learn a better model  $M_{i+1}$
- Key challenge?

## Model-based RL Example

- Say world is deterministic
  - and no wind
- Lets say the agent first discovers the path to bad reward first
- Will the agent ever learn the optimal policy?
  - won't have any information about some states or state-action pairs

	1	2	3	4
A	↓	?	?	+100
B	↓	?	?	?
C	→	→	→	-2

# Model-based RL

- Learn an initial model  $M_0$
- Loop
  - VI/PI on  $M_i$  to compute policy  $\pi_i$
  - Execute  $\pi_i$  to generate data
  - Learn a better model  $M_{i+1}$
- Key challenge
  - Just executing  $\pi_i$  is not enough!
  - It may miss important regions
  - Needs to explore new regions

## TD Learning $\rightarrow$ TD ( $V^*$ ) Learning

- Can we do TD-like updates on  $V^*$ ?
- $V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$
- Hmm... what to do?
  - RHS should be expectation.
  - Instead of  $V^*$  write all equations in  $Q^*$



## Bellman Equations ( $V^*$ ) $\rightarrow$ Bellman Equations ( $Q^*$ )

- $V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$
- $Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$
- $Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma \max_{a'} Q^*(s', a')]$
- VI  $\rightarrow$  Q-Value Iteration
- TD Learning  $\rightarrow$  Q Learning

# Q Learning

- Directly learn  $Q^*(s,a)$  values
- Receive a sample  $(s, a, s', r)$
- Your old estimate  $Q(s,a)$
- New sample value:  $r + \gamma \max_{a'} Q(s', a')$

Nudge the estimates:

- $Q(s,a) \leftarrow Q(s,a) + \alpha(r + \gamma \max_{a'} Q(s', a') - Q(s,a))$
- $Q(s,a) \leftarrow (1 - \alpha)Q(s,a) + \alpha(r + \gamma \max_{a'} Q(s', a'))$

# Q Learning Algorithm

- **For all  $s, a$** 
  - Initialize  $Q(s, a) = 0$
- **Repeat Forever**
  - Where are you?  $s$ .
  - Choose some action  $a$
  - Execute it in real world:  $(s, a, r, s')$
  - Do update:
$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha(r + \gamma \max_{a'} Q(s', a'))$$

Is an *off policy learning* algorithm

# Properties

- Q Learning converges to optimal values  $Q^*$ 
  - Irrespective of initialization,
  - Irrespective of action choice policy
  - Irrespective of learning rate
- as long as
  - states/actions finite, all rewards bounded
  - No (s,a) is starved: infinite visits over infinite samples
  - Learning rate decays with visits to state-action pairs
    - but not too fast decay. ( $\sum_i \alpha(s,a,i) = \infty$ ,  $\sum_i \alpha^2(s,a,i) < \infty$ )

# Q Learning Algorithm

- For all  $s, a$ 
  - Initialize  $Q(s, a) = 0$
- Repeat Forever
  - Where are you?  $s$ .
  - Choose some action  $a$
  - Execute it in real world:  $(s, a, r, s')$
  - Do update:
$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha(r + \gamma \max_{a'} Q(s', a'))$$

How to choose?

new: exploration

greedy: exploitation

# Exploration vs. Exploitation Tradeoff

- A fundamental tradeoff in RL
- **Exploration:** must take actions that may be suboptimal but help discover new rewards and in the long run increase utility
- **Exploitation:** must take actions that are known to be good (and seem currently optimal) to optimize the overall utility
- Slowly move from exploration → exploitation

# Explore/Exploit Policies

- Simplest scheme:  $\epsilon$ -greedy
  - Every time step flip a coin
  - With probability  $1-\epsilon$ , take the greedy action
  - With probability  $\epsilon$ , take a random action
- Problem
  - Exploration probability is constant
- Solutions
  - Lower  $\epsilon$  over time
  - Use an exploration function

# Explore/Exploit Policies

- Boltzmann Exploration

- Select action  $a$  with probability

- $$\Pr(a|s) = \frac{\exp(Q(s,a)/T)}{\sum_{a' \in A} \exp(Q(s,a')/T)}$$

- T: Temperature

- Similar to simulated annealing
- Large T: uniform, Small T: greedy
- Start with large T and decrease with time

- GLIE: greedy in the limit of infinite exploration



# Explore/Exploit Policies

- Exploration Functions
  - stop exploring actions whose badness is established
  - continue exploring other actions
- Let  $Q(s,a) = q$ ,  $\#visits(s,a) = n$
- E.g.:  $f(q, n) = q + k/n$ 
  - Unexplored states have infinite  $f$
  - Highly explored bad states have low  $f$
- Modified Q update
  - $Q(s,a) \leftarrow (1 - \alpha)Q(s,a) + \alpha(r + \gamma \max_{a'} f(Q(s', a'), N(s', a')))$

States leading to unexplored states are also preferred<sup>41</sup>

# Explore/Exploit Policies

- A Famous Exploration Policy: UCB
  - Upper Confidence Bound

$$\pi_{UCT}(s) = \arg \max_a Q(s, a) + c \sqrt{\frac{\ln n(s)}{n(s, a)}}$$

## Value Term:

favors actions that looked good historically

## Exploration Term:

actions get an exploration bonus that grows with  $\ln(n)$

Optimistic in the Face of Uncertainty

# Model based vs. Model Free RL

## ■ Model based

- estimate  $O(|\mathcal{S}|^2|\mathcal{A}|)$  parameters
- requires relatively larger data for learning
- can make use of background knowledge easily

## ■ Model free

- estimate  $O(|\mathcal{S}||\mathcal{A}|)$  parameters
- requires relatively less data for learning

# Generalizing Across States

- Basic Q-Learning (or VI) keeps a table of all q-values
- In realistic situations, we cannot possibly learn about every single state!
  - Too many states to visit them all in training
  - Too many states to hold the q-tables in memory
- Instead, we want to generalize:
  - Learn about some small number of training states from experience
  - Generalize that experience to new, similar situations
  - This is a fundamental idea in machine learning

# Feature-based Representation

- Describe a state using vector of features
- We can write a q function using a few weights:

$$Q(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \dots + w_n f_n(s, a)$$

- **Advantage:** our experience is summed up in a few powerful numbers ( $w_i$ )
- **Disadvantage:** states may share features but actually be very different in value!

# Approximate Q-Learning

$$Q(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \dots + w_n f_n(s, a)$$

- Exact Q-Learning

- $Q(s, a) \leftarrow Q(s, a) + \alpha (r + \gamma \max_{a'} Q(s', a') - Q(s, a))$

difference


- Q-Learning with linear function approximation

- $w_m \leftarrow w_m + \alpha (r + \gamma \max_{a'} Q(s', a') - Q(s, a)) f_m(s, a)$

- Move feature weights up/down based on difference and feature values

# Optimization: Least Squares

$$\text{total error} = \sum_i (y_i - \hat{y}_i)^2 = \sum_i \left( y_i - \sum_k w_k f_k(x_i) \right)^2$$



# Minimizing Error

Imagine we had only one point  $x$ , with features  $f(x)$ , target value  $y$ , and weights  $w$ :

$$\begin{aligned}\text{error}(w) &= \frac{1}{2} \left( y - \sum_k w_k f_k(x) \right)^2 \\ \frac{\partial \text{error}(w)}{\partial w_m} &= - \left( y - \sum_k w_k f_k(x) \right) f_m(x) \\ w_m &\leftarrow w_m + \alpha \left( y - \sum_k w_k f_k(x) \right) f_m(x)\end{aligned}$$

Approximate q update


explained:  $w_m \leftarrow w_m + \alpha \left[ r + \gamma \max_a Q(s', a') - Q(s, a) \right] f_m(s, a)$

“target”

“prediction”



# Overfitting and Limited Capacity Approximations



Low capacity generalizes better

Issue: linear approximation not powerful enough in practice

Deep Learning!

## Summary: RL

RL is a very general AI problem  
most general single agent?

Main idea:  $\text{expectation}_P$  as avg of samples  
sampling distribution is  $P$

Agent learns as it gathers experience

Exploration-exploitation tradeoff

Function approximation is key: deep RL is the rage!

# Applications

- Stochastic Games
- Robotics: navigation, helicopter maneuvers...
- Finance: options, investments
- Communication Networks
- Medicine: Radiation planning for cancer
- Controlling workflows
- Optimize bidding decisions in auctions
- Traffic flow optimization
- Aircraft queueing for landing; airline meal provisioning
- Optimizing software on mobiles
- Forest firefighting
- ...