

Maintaining user/server state: cookies

Web sites and client browser use

cookies to maintain some state
between transactions

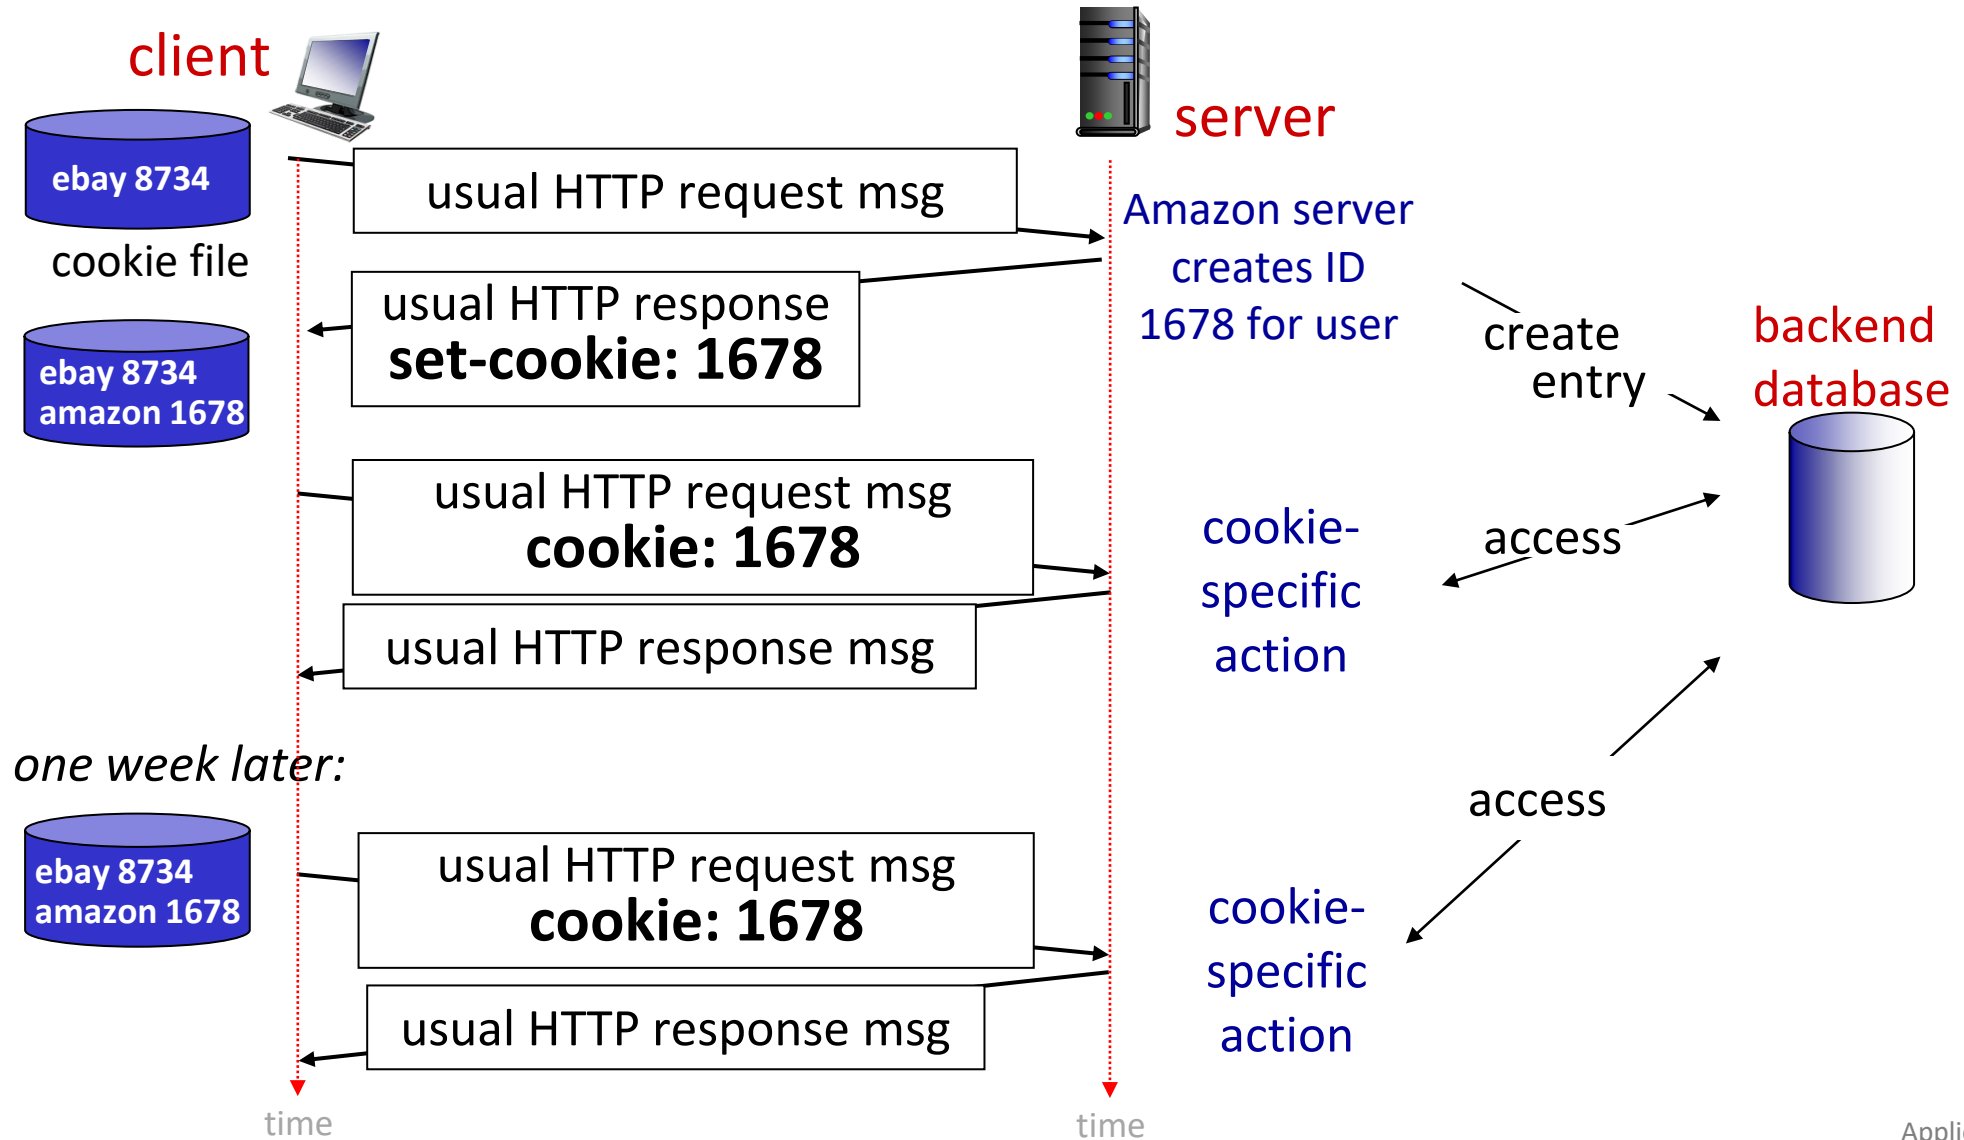
four components:

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in next HTTP *request* message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

Example:

- Susan uses browser on laptop, visits specific e-commerce site for first time
- when initial HTTP requests arrives at site, site creates:
 - unique ID (aka “cookie”)
 - entry in backend database for ID
- subsequent HTTP requests from Susan to this site will contain cookie ID value, allowing site to “identify” Susan

Maintaining user/server state: cookies



HTTP cookies: comments

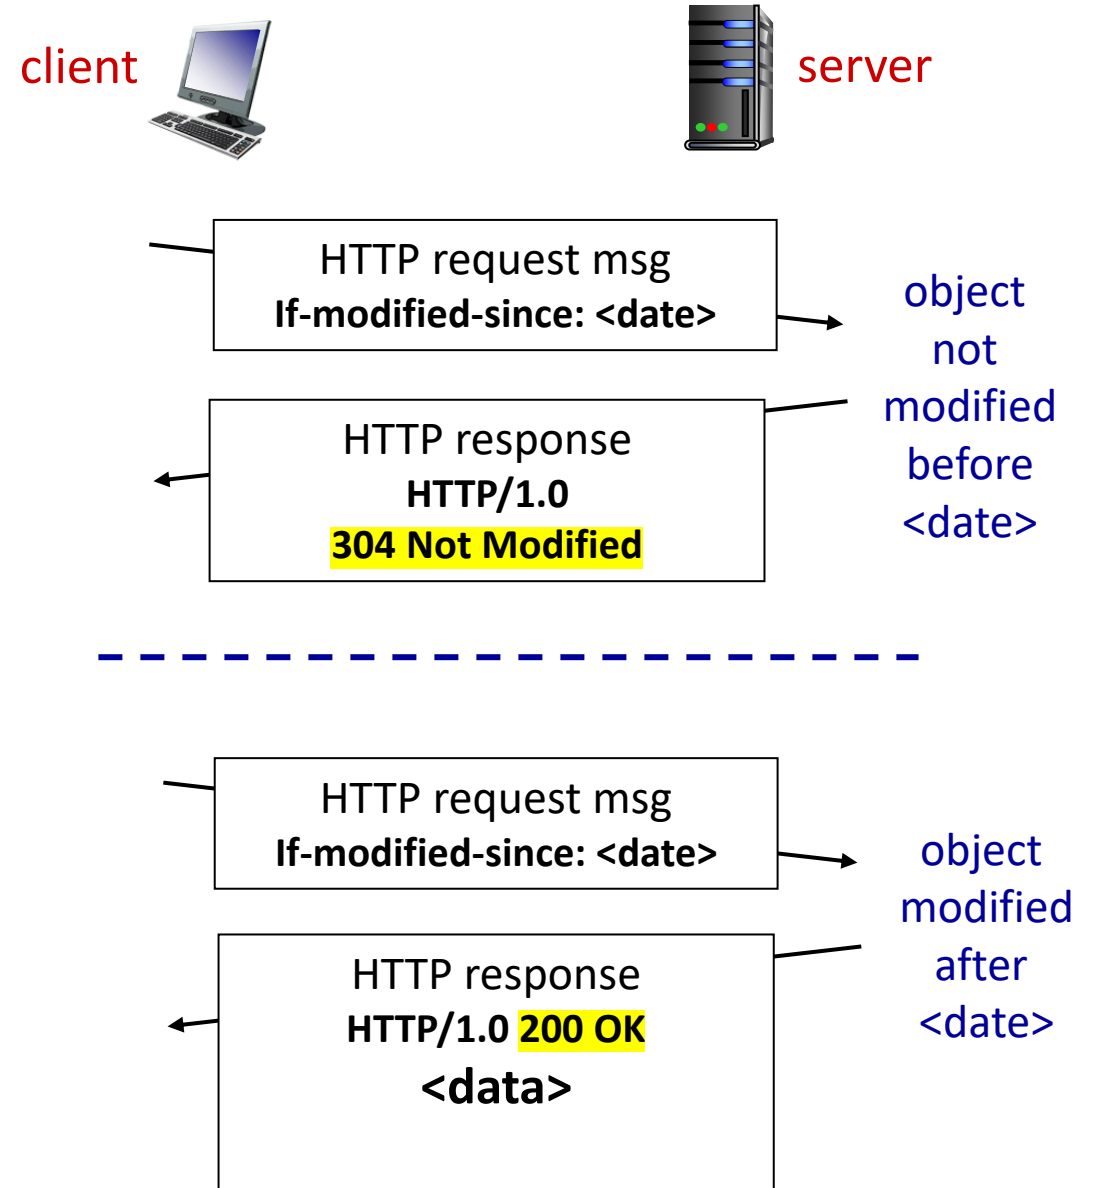
What cookies can be used for:

- shopping carts
- recommendations
- user session state

Conditional GET

Goal: don't send object if cache has up-to-date cached version

- no object transmission delay
- lower link utilization
- **cache:** specify date of cached copy in HTTP request
If-modified-since: <date>
- **server:** response contains no object if cached copy is up-to-date:
HTTP/1.0 304 Not Modified



HTTP/2

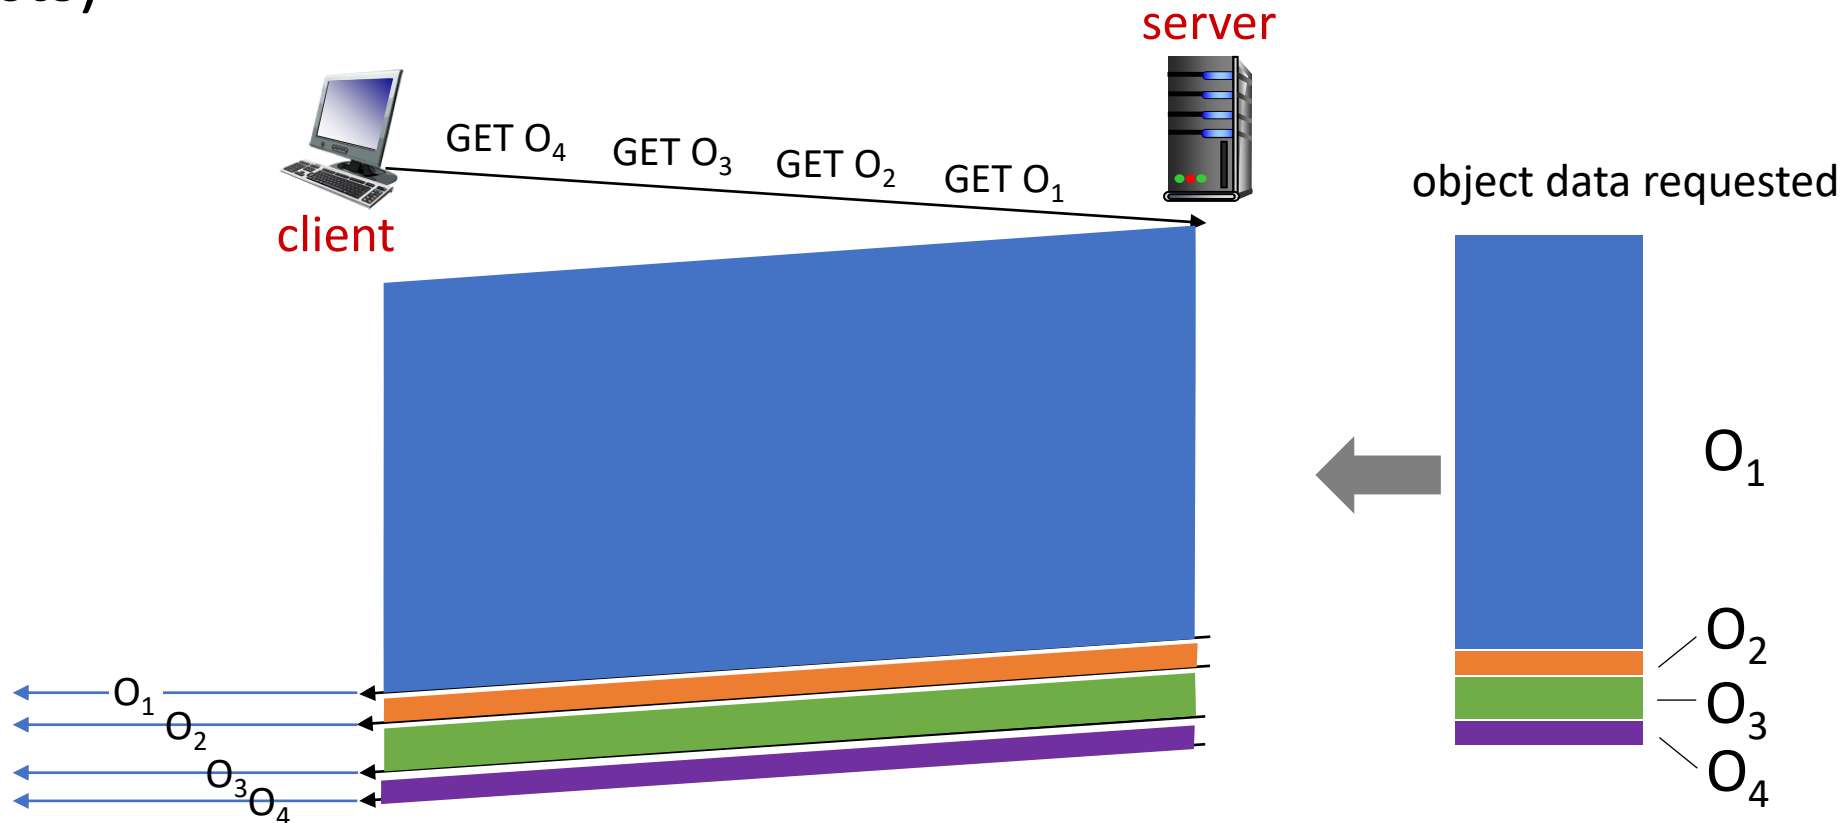
Key goal: decreased delay in multi-object HTTP requests

HTTP1.1: introduced multiple, pipelined GETs over single TCP connection

- server responds *in-order* (FCFS: first-come-first-served scheduling) to GET requests
- with FCFS, small object may have to wait for transmission (head-of-line (HOL) blocking) behind large object(s)
- loss recovery (retransmitting lost TCP segments) stalls object transmission

HTTP/2: mitigating HOL blocking

HTTP 1.1: client requests 1 large object (e.g., video file, and 3 smaller objects)



objects delivered in order requested: O_2 , O_3 , O_4 wait behind O_1

HTTP/2

Key goal: decreased delay in multi-object HTTP requests

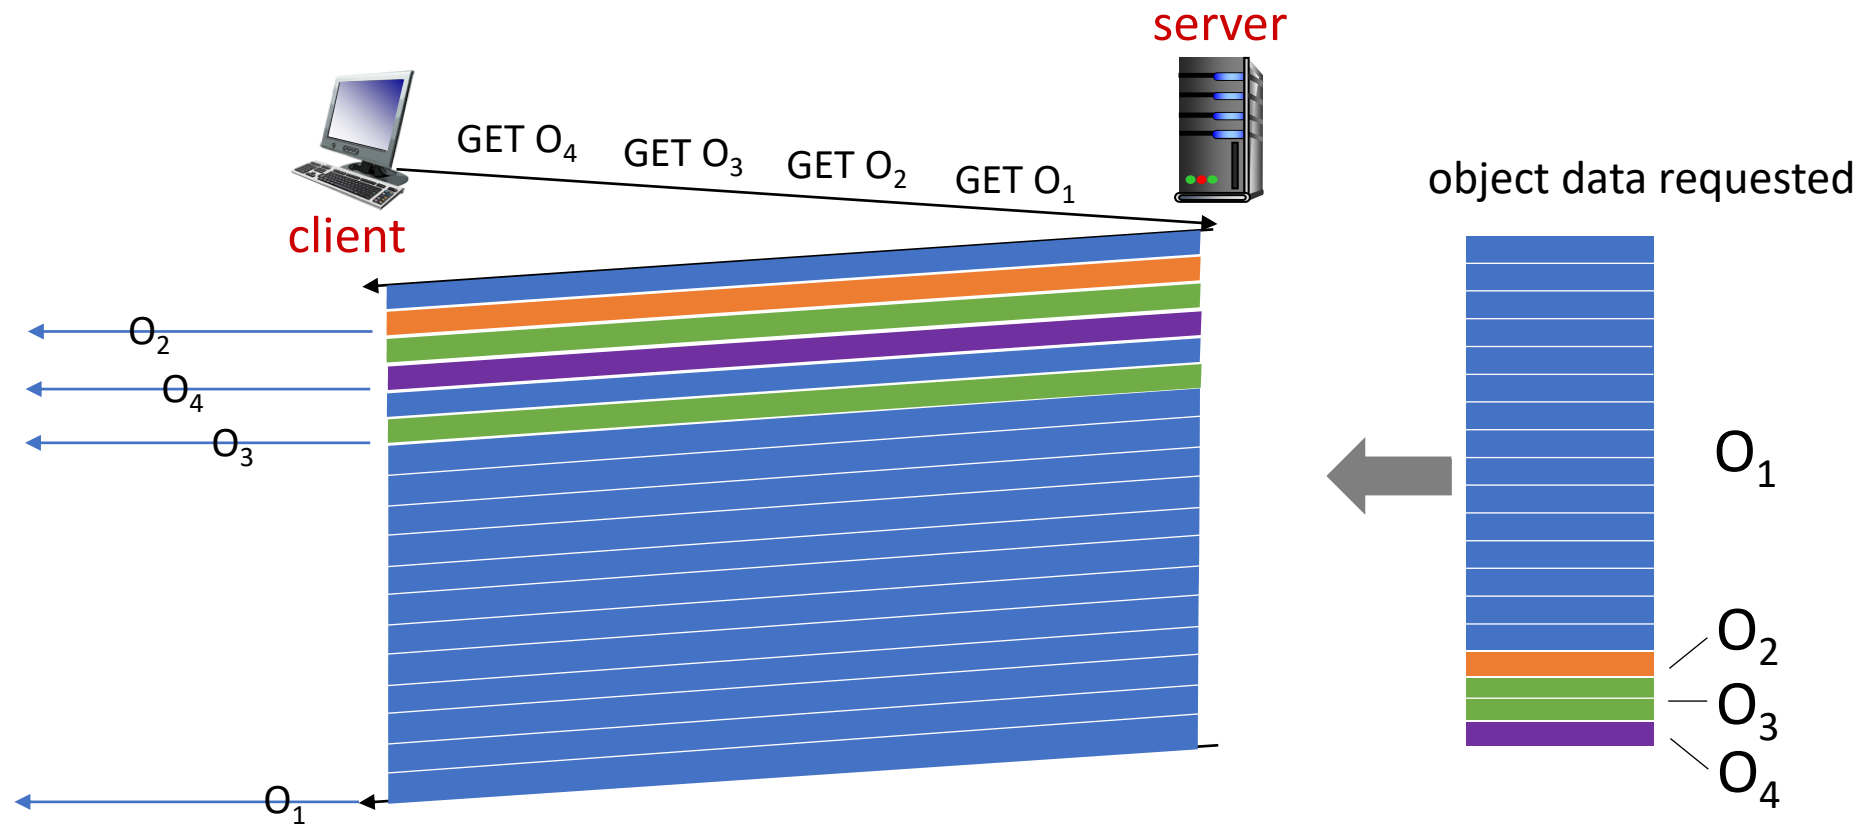
HTTP/2: [RFC 7540, 2015] **increased flexibility** at *server* in sending objects to client:

- methods, status codes, most header fields **unchanged** from HTTP 1.1
- **transmission order** of requested objects based on **client-specified object priority** (not necessarily FCFS)
- **divide objects into frames**, schedule frames to mitigate **HOL blocking**

HTTP/3: adds **security**, per object **error- and congestion-control** (more pipelining) over UDP [**Upcoming major version of HTTP**]

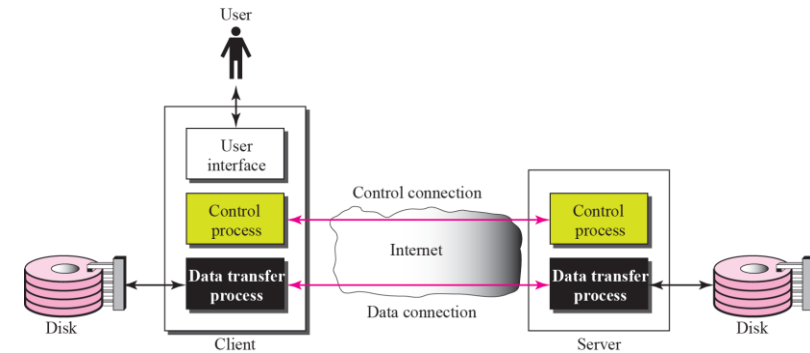
HTTP/2: mitigating HOL blocking

HTTP/2: objects **divided into frames**, frame transmission interleaved



O_2 , O_3 , O_4 delivered quickly, O_1 slightly delayed

File Transfer Protocol (FTP)



- It is like HTTP.
- You connect to a FTP server and can "get" or "put" files, after which data transfer happens.
- Some differences compared to HTTP:
 - (1) FTP sends **control [TCP,21]** and **data [TCP,20]** through separate TCP connections.
 - Control connection: user **identification, password**, commands to change remote directory, **commands** to put and get files.
 - **Data** connection: It is used to **actually send** the file.
 - **FTP is said to send control information out of band unlike HTTP.**
 - **Control connection remains open** throughout the duration of user session, but a new data connection is created for each file transferred within a session. (i.e., **data connection are non-persistent**)
- (2) **FTP keeps state** over a session, e.g., user's **current working directory** etc.

Application Layer: Overview

- Principles of network applications
- Web and HTTP
- The Domain Name System DNS
- E-mail, SMTP, IMAP
- P2P applications
- video streaming and content distribution networks
- socket programming with UDP and TCP



DNS: Domain Name System

people: many identifiers:

- name, passport #, Aadhar #

Internet hosts, routers:

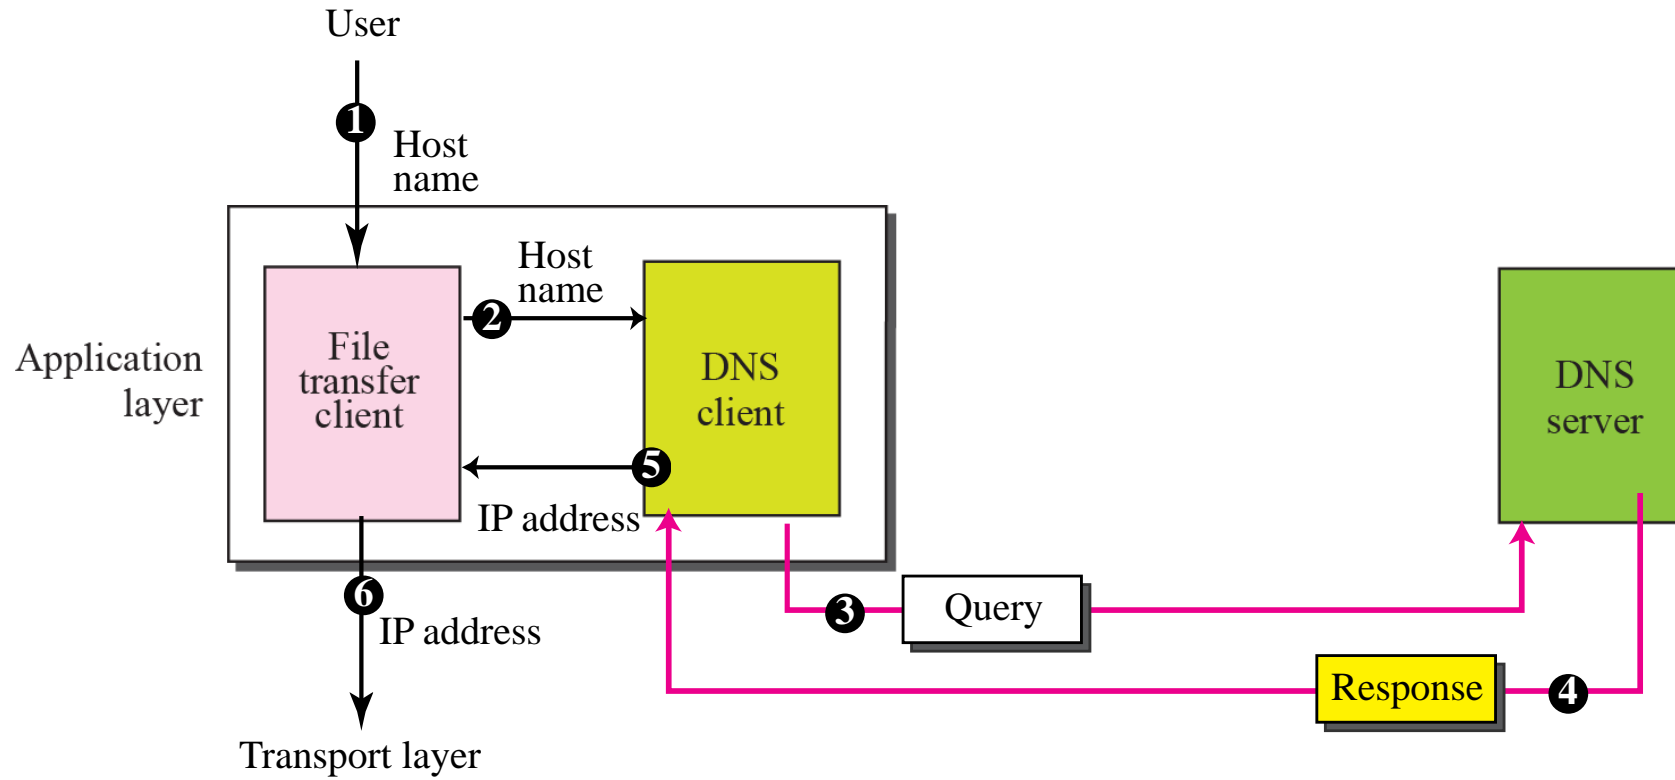
- IP address (32 bit) - used for addressing datagrams
- “name”, e.g., cs.umass.edu - used by humans

Q: how to map between IP address and name, and vice versa ?

Domain Name System:

- *distributed database* implemented in hierarchy of many *name servers*
- *application-layer protocol:* hosts, name servers communicate to resolve names (address/name translation)
 - note: core Internet function, implemented as *application-layer protocol*

Purpose of DNS



DNS: services, structure

DNS services

- hostname to IP address translation
- load distribution
 - replicated Web servers: many IP addresses correspond to one name


Q: Why not centralize DNS?

- single point of failure
- traffic volume
- distant centralized database
- maintenance

A: doesn't scale!

- Comcast DNS servers alone: 600B DNS queries per day

DNS: a distributed, hierarchical database



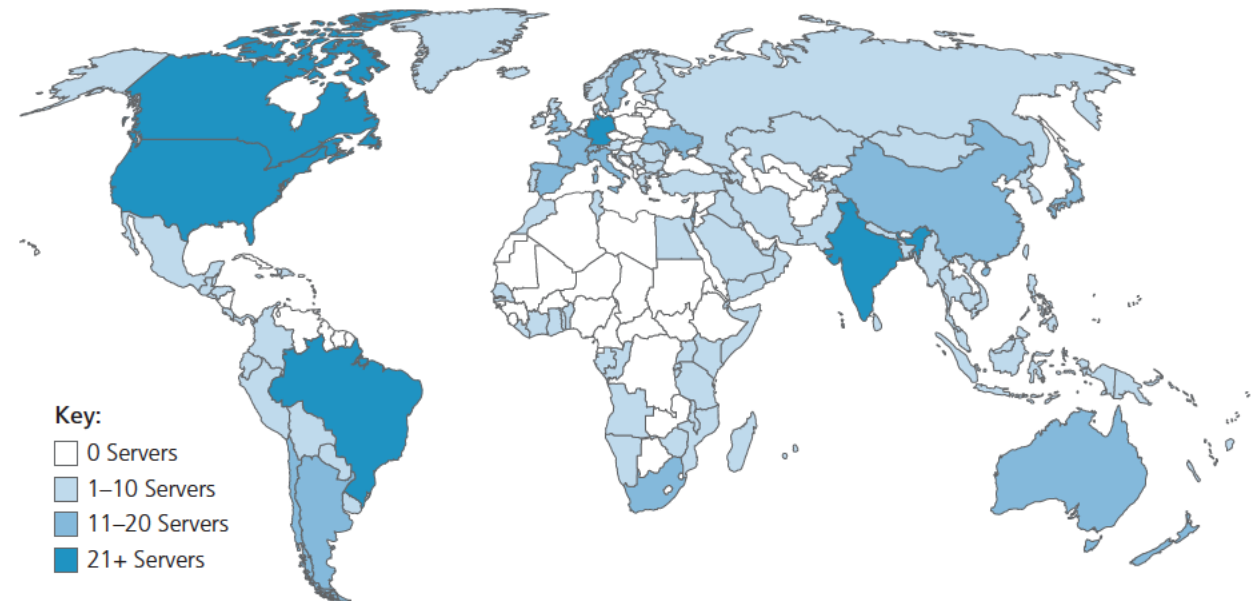
Client wants IP address for `www.amazon.com`; 1st approximation:

- client queries root server to find `.com` DNS server
- client queries `.com` DNS server to get `amazon.com` DNS server
- client queries `amazon.com` DNS server to get IP address for `www.amazon.com`

DNS: root name servers

- official, contact-of-last-resort by name servers that can not resolve name
- *incredibly important* Internet function
 - Internet couldn't function without it!
 - DNSSEC – provides security (authentication and message integrity)
- ICANN (Internet Corporation for Assigned Names and Numbers) manages root DNS domain

13 logical root name “servers”
worldwide each “server” replicated
many times (~200 servers in US)



TLD: authoritative servers

Top-Level Domain (TLD) servers:

- Generic top-level domains (gTLD): .com, .org, .net, .edu, .aero, .jobs, .net, .edu
- Country-code top-level domains (ccTLD): all country domains, e.g.: .in .cn, .uk, .fr, .ca, .jp

Authoritative DNS servers:

- organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
- can be maintained by organization or service provider

Local DNS name servers


- does not strictly belong to hierarchy
- each ISP (residential ISP, company, university) has one
 - also called “default name server”
- when host makes DNS query, query is sent to its local DNS server
 - has local cache of recent name-to-address translation pairs (but may be out of date!)
 - acts as proxy, forwards query into hierarchy
 - The local DNS servers are statically configured with the identity of the root servers.

DNS name resolution: iterated query

Example: host at `engineering.nyu.edu` wants IP address for `gaia.cs.umass.edu`

Iterated query:

- contacted server replies with **name of server to contact**
- “I don’t know this name, but ask this server”



DNS name resolution: recursive query

Example: host at `engineering.nyu.edu` wants IP address for `gaia.cs.umass.edu`

Recursive query:

- puts **burden** of name resolution on **contacted name server**
- heavy load at upper levels of hierarchy?

