

# Transport Layer



Anand Baswade

[anand@iitbhilai.ac.in](mailto:anand@iitbhilai.ac.in)

# TCP Sender (simplified)

event: data received from application

- create segment with seq #
- seq # is byte-stream number of first data byte in segment
- start timer if not already running
  - think of timer as for oldest unACKed segment
  - expiration interval: **TimeoutInterval**


event: timeout

- retransmit segment that caused timeout
- restart timer

event: ACK received

- if ACK acknowledges previously unACKed segments
  - update what is known to be ACKed
  - start timer if there are still unACKed segments

# TCP Receiver: ACK generation [RFC 5681]



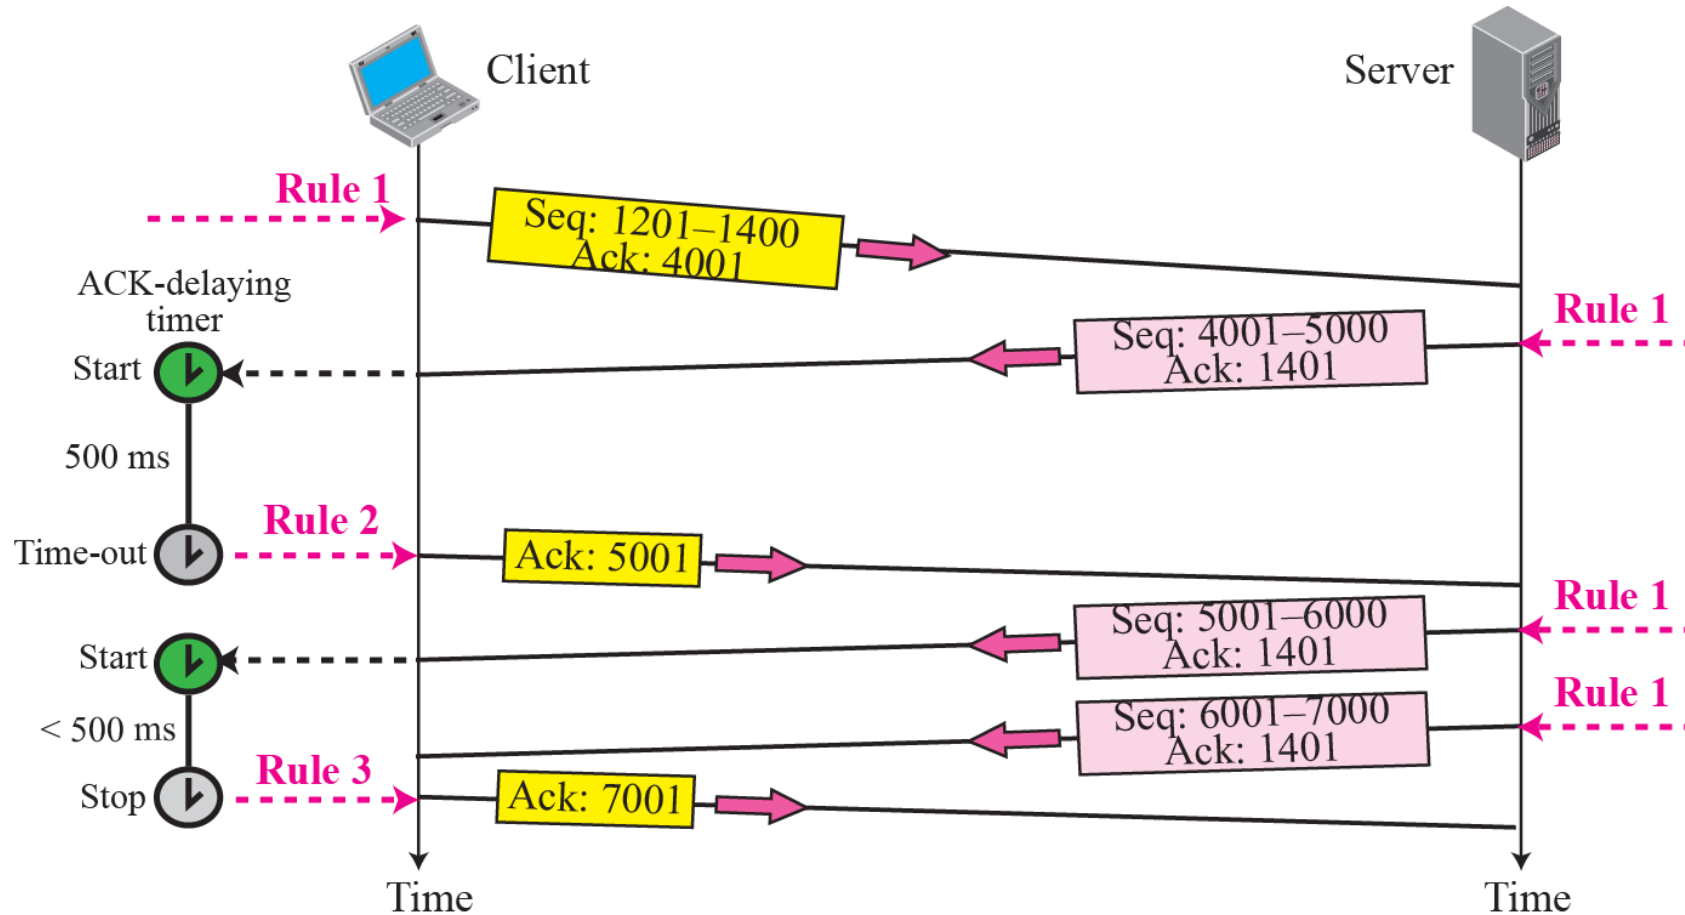
# Rules for Generating the ACKs

1. When one end sends a data segment to the other end, it must include an ACK. That gives the next sequence number it expects to receive. (Piggyback)
2. The receiver needs to delay sending (until another segment arrives or 500ms) an ACK segment if there is only one outstanding in-order segment. It prevents ACK segments from creating extra traffic.
3. There should not be more than 2 in-order unacknowledged segments at any time. It prevents the unnecessary retransmission


# Rules for Generating the ACKs Cont..

4. When a segment arrives with an out-of-order sequence number that is higher than expected, the receiver immediately sends an ACK segment announcing the sequence number of the next expected segment. (for fast retransmission)
5. When a missing segment arrives, the receiver sends an ACK segment to announce the next sequence number expected.
6. If a duplicate segment arrives, the receiver immediately sends an ACK.

# *Some Scenarios: Normal operation*



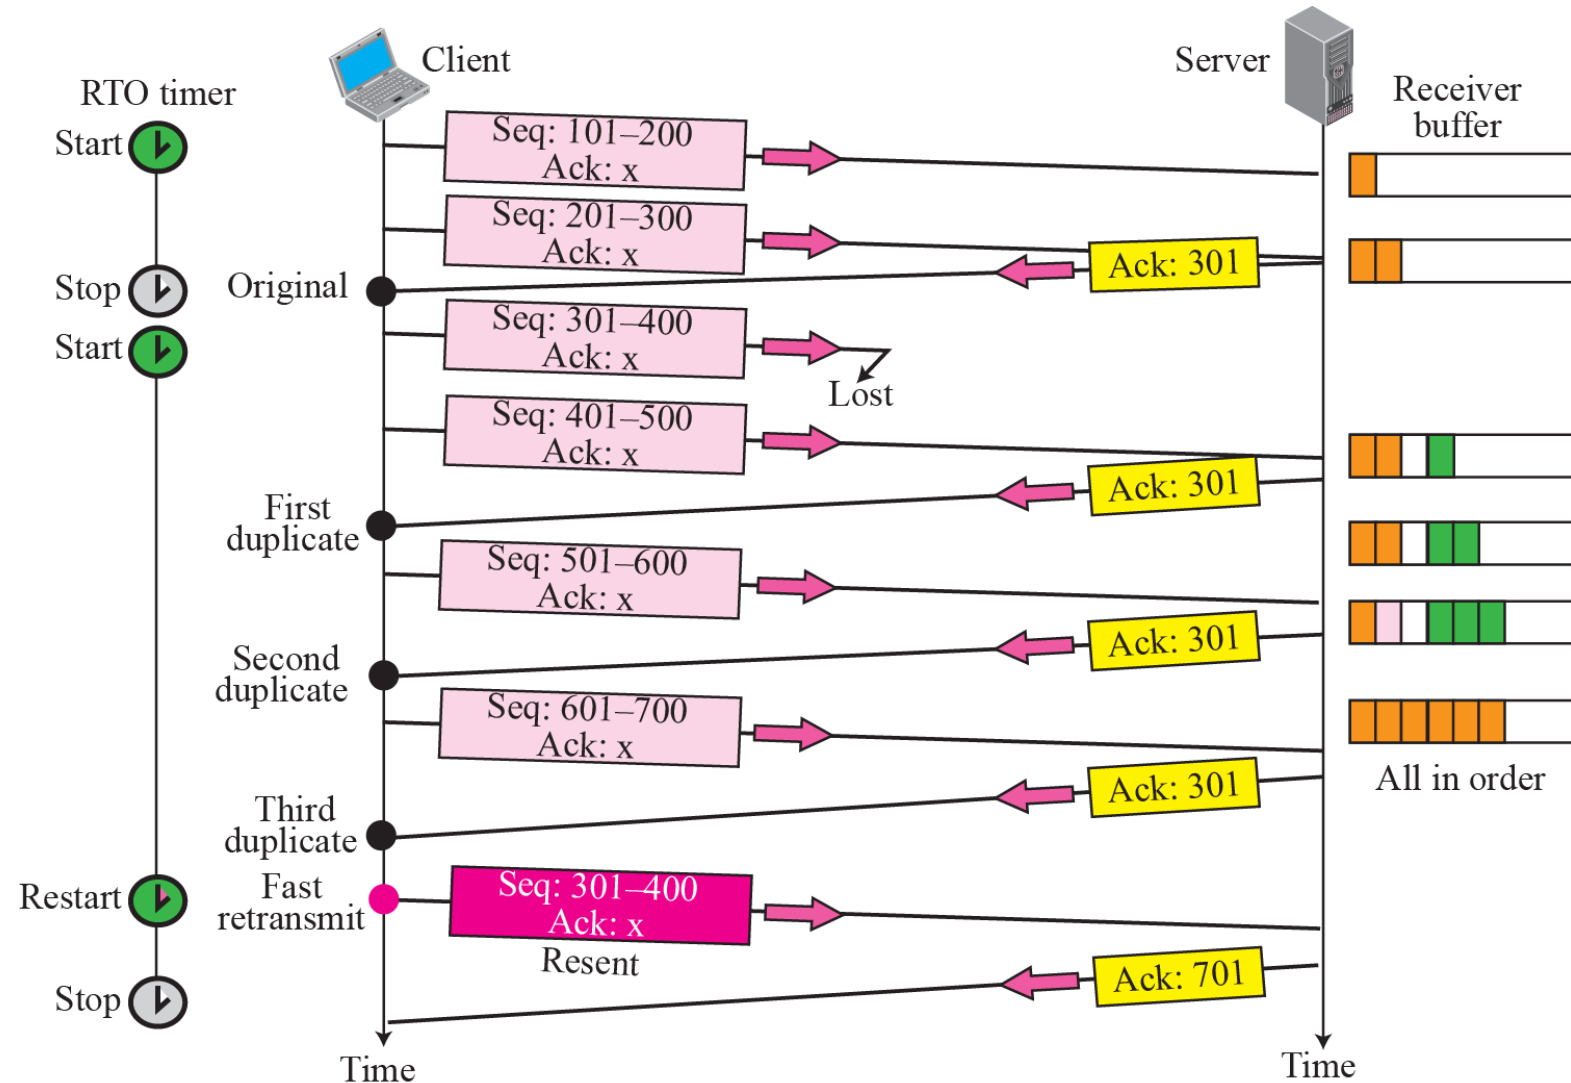
# Lost segment



*The receiver TCP delivers only ordered data to the process.*


# Fast retransmission

Receipt of three  
duplicate ACKs  
indicates 3  
segments received  
after a missing  
segment – lost  
segment is likely. So  
retransmit!






# Lost acknowledgment



# *Lost acknowledgment corrected by resending a segment*



# ACK and Out of Order Handling in TCP

## Acknowledgement in TCP – Cumulative acknowledgement

Receiver has received bytes 0, 1, 2, \_, 4, 5, 6, 7

- TCP sends a cumulative acknowledgement with ACK number 3, acknowledging everything up to byte 2
- Once 4 is received, a duplicate ACK with ACK number 3 (next expected byte) is forwarded
- After timeout, sender retransmits byte 3
- Once byte 3 is received, it can send another cumulative ACK with ACK number 8 (next expected byte)

# TCP round trip time, timeout

Q: how to set TCP timeout value?

- longer than RTT, but RTT varies!
- *too short*: premature timeout, unnecessary retransmissions
- *too long*: slow reaction to segment loss


Q: how to estimate RTT?

- SampleRTT: measured time from segment transmission until ACK receipt
  - ignore retransmissions
- SampleRTT will vary, want estimated RTT “smoother”
  - average several *recent* measurements, not just current SampleRTT

# TCP round trip time, timeout

$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- exponential weighted moving average (EWMA)
- influence of past sample decreases exponentially fast
- typical value:  $\alpha = 0.125$



# TCP round trip time, timeout

- timeout interval: **EstimatedRTT** plus “safety margin”
  - large variation in **EstimatedRTT**: want a larger safety margin

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$



↑  
estimated RTT

↑  
“safety margin”


- **DevRTT**: EWMA of **SampleRTT** deviation from **EstimatedRTT**:

$$\text{DevRTT} = (1 - \beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically,  $\beta = 0.25$ )


# TCP flow control

Q: What happens if **network layer delivers** data faster than **application layer** removes data from socket buffers?



# TCP flow control


Q: What happens if network layer delivers data faster than application layer removes data from socket buffers?





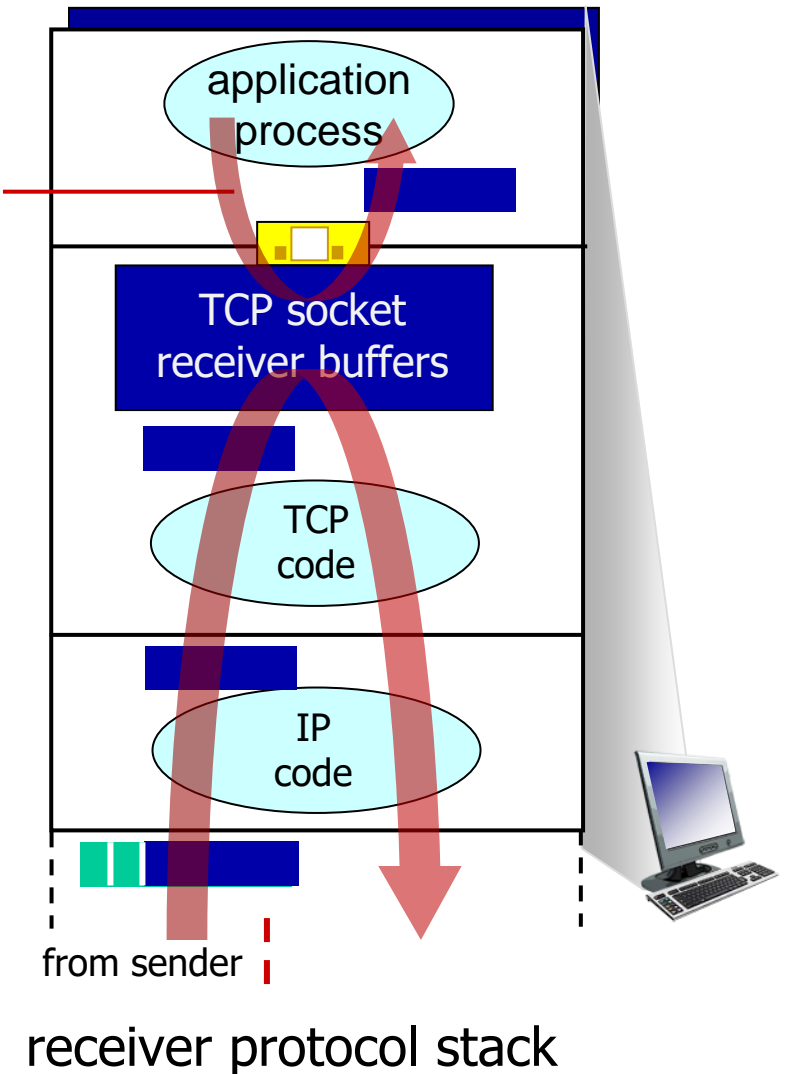
# TCP flow control

Q: What happens if network layer delivers data faster than application layer removes data from socket buffers?



flow control: # bytes  
receiver willing to accept

Application removing  
data from TCP socket  
buffers

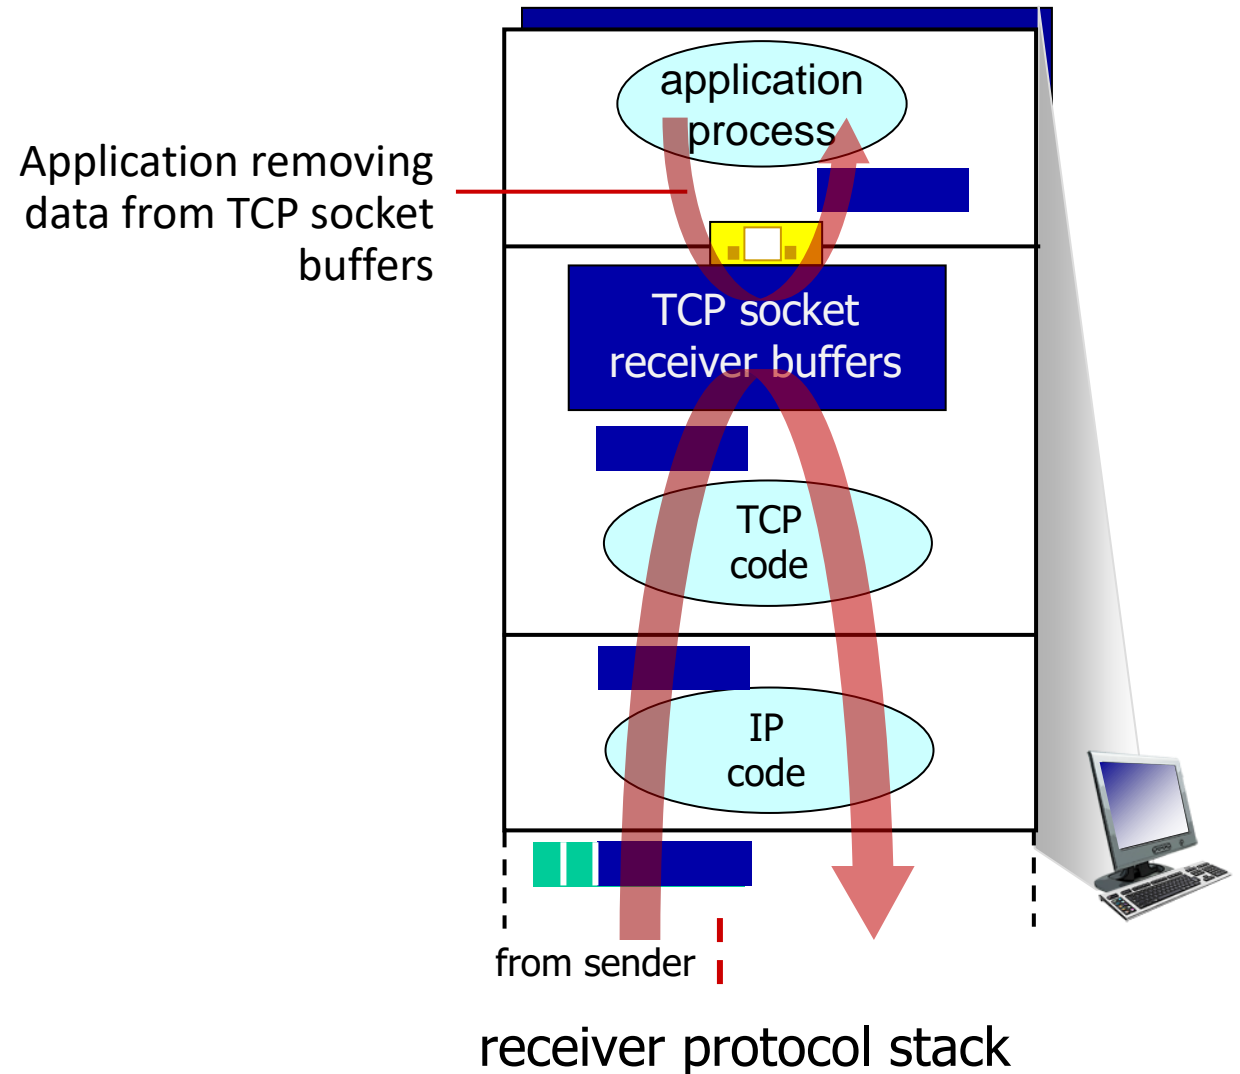


# TCP flow control

Q: What happens if network layer delivers data faster than application layer removes data from socket buffers?


## —flow control—

receiver controls sender, so sender won't overflow receiver's buffer by transmitting too much, too fast



# TCP flow control

- TCP receiver “advertises” free buffer space in **rwnd** field in TCP header
  - **RcvBuffer** size set via socket options (typical default is 4096 bytes)
  - many operating systems autoadjust **RcvBuffer**
- guarantees receive buffer will not overflow




TCP receiver-side buffering

# TCP flow control

- TCP receiver “advertises” free buffer space in **rwnd** field in TCP header
  - **RcvBuffer** size set via socket options (typical default is 4096 bytes)
  - many operating systems autoadjust **RcvBuffer**
- sender **limits amount** of unACKed (“in-flight”) data to **received rwnd**
- guarantees receive buffer will not overflow

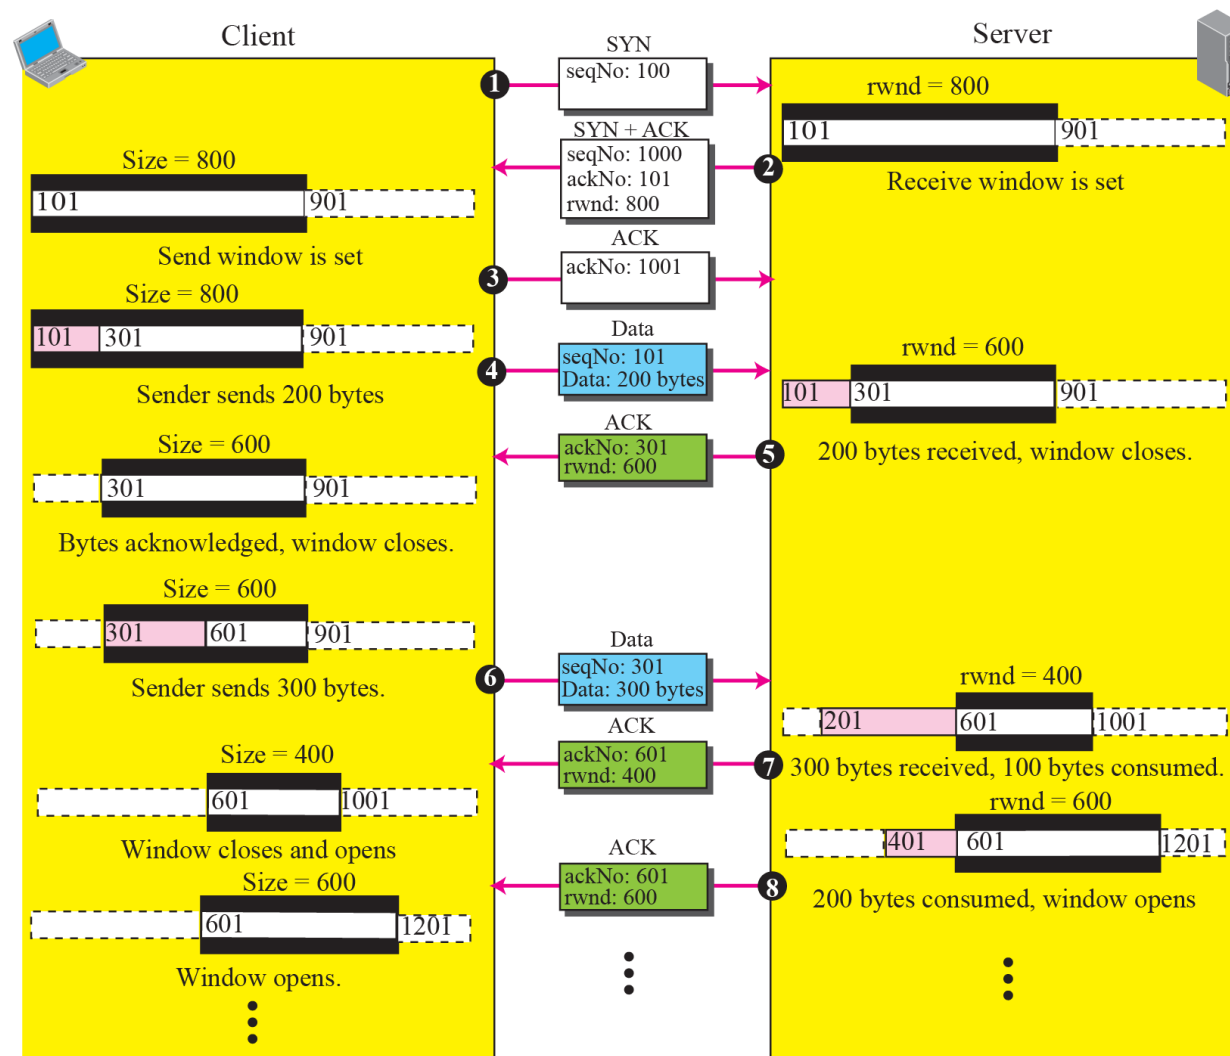
flow control: # bytes receiver willing to accept



TCP segment format

# An example of flow control

**Note:** We assume only unidirectional communication from client to server. Therefore, only one window at each side is shown.



# Principles of congestion control

## Congestion:

- informally: “too many sources sending too much data **too fast for network** to handle”
  - **long delays** (queueing in router buffers)
  - **packet loss** (buffer **overflow** at routers)
- different from **flow control**!
- a top-10 problem!



**congestion control:**

too many senders,  
sending too fast




**flow control:** one sender  
too fast for one receiver

# Approaches towards congestion control

## End-end congestion control:


- no explicit feedback from network
- congestion *inferred* from observed loss, delay
- approach taken by TCP



# Approaches towards congestion control

## Network-assisted congestion control:

- routers provide *direct* feedback to sending/receiving hosts with flows passing through congested router
- may indicate congestion level or explicitly set sending rate
- TCP ECN (Explicit Congestion Notification), ATM (Asynchronous Transfer Mode)





# Chapter 3: roadmap

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- Principles of congestion control
- **TCP congestion control**
- Evolution of transport-layer functionality



# TCP: Triggering congestion control

- Two ways to trigger a congestion notification in TCP – (1) RTO, (2) Duplicate ACK
- RTO: A sure indication of congestion, however time consuming
- Duplicate ACK: Receiver sends a duplicate ACK when it receives out of order segment
  - A loose way of indicating congestion
  - TCP arbitrarily assumes that THREE duplicate ACKs (DUPACKs) imply that a packet has been lost – triggers congestion control mechanism
  - Retransmit the lost packet and trigger congestion control

# TCP congestion control: AIMD


- *approach*: senders can increase sending rate until packet loss (congestion) occurs, then decrease sending rate on loss event.

## Additive Increase

increase sending rate by 1 maximum segment size every RTT until loss detected

## Multiplicative Decrease

cut sending rate in half at each loss event



- Chiu and Jain (1989): Let  $w(t)$  be the sending rate.  $a$  ( $a > 0$ ) is the additive increase factor, and  $b$  ( $0 < b < 1$ ) is the multiplicative decrease factor

$$w(t+1) = \begin{cases} w(t) + a & \text{if congestion is not detected} \\ w(t) \times b & \text{if congestion is detected} \end{cases}$$