

# Transport Layer




Anand Baswade

[anand@iitbhilai.ac.in](mailto:anand@iitbhilai.ac.in)

# TCP fairness

**Fairness goal:** if  $K$  TCP sessions share same bottleneck link of bandwidth  $R$ , each should have average rate of  $R/K$




## Goal:

1. Fairness: Every user should get a fair share of link
2. Capacity: Network link bandwidth should be well utilized

# Q: is TCP Fair?

Example: two competing TCP sessions:

- additive increase – Increases Efficiency
- multiplicative decrease decreases throughput proportionally - Increases Fairness



*Is TCP fair?*

**A:** Yes, under idealized assumptions:

- same RTT
- fixed number of sessions

# Fairness: must all network apps be “fair”?

## Fairness and UDP

- multimedia apps often do not use TCP
  - do not want rate throttled by congestion control
- instead use UDP:
  - send audio/video at constant rate, tolerate packet loss
- there is no “Internet police” policing use of congestion control


## Fairness, parallel TCP connections

- application can open *multiple* parallel connections between two hosts
- web browsers do this , e.g., link of rate  $R$  with 9 existing connections:
  - new app asks for 1 TCP, gets rate  $R/10$
  - new app asks for 11 TCPs, gets  $R/2$

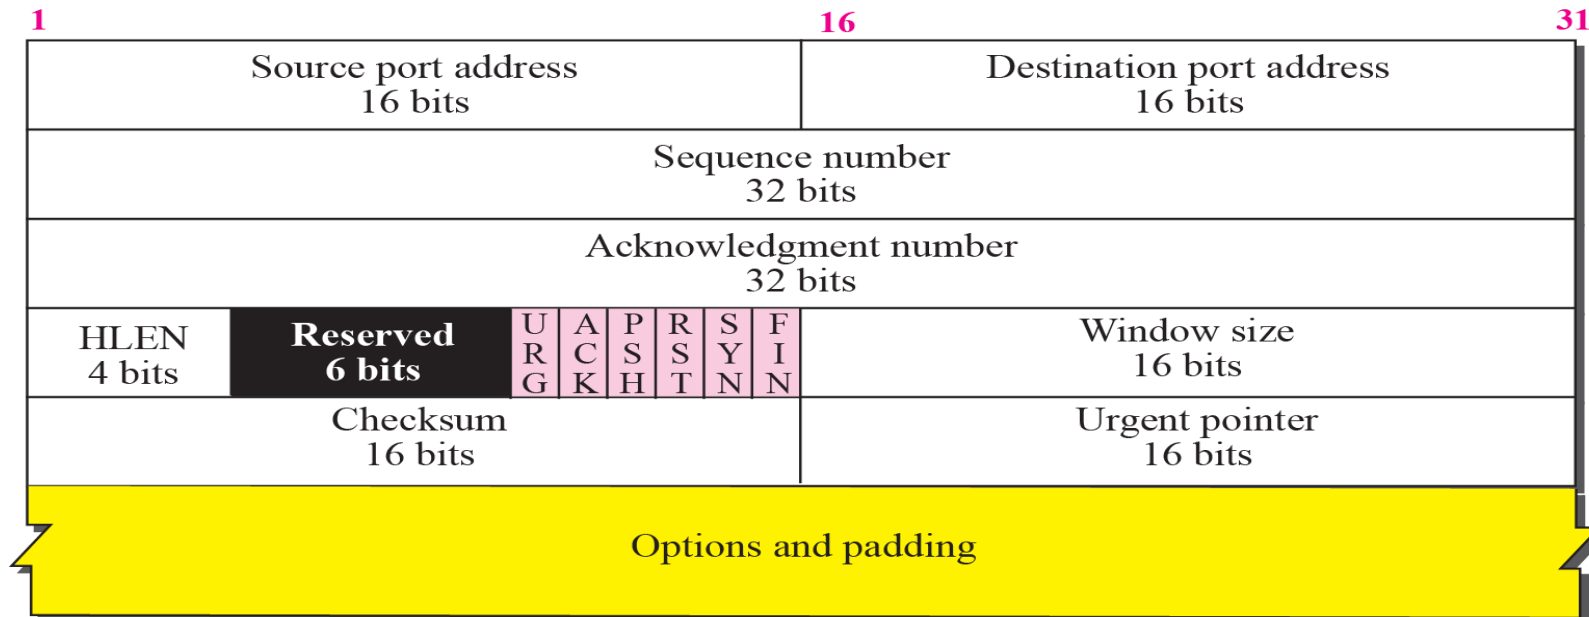
# TCP Options

- The TCP header can have up to 40 bytes of optional information. Options convey additional information to the destination.
- We can define two categories of options: 1-byte options and multiple-byte options.
- The first category contains two types of options: end of option list and no operation.
- The second category, in most implementations, contains five types of options: maximum segment size, window scale factor, timestamp, SACK-permitted, and SACK.

# *TCP segment with options format*

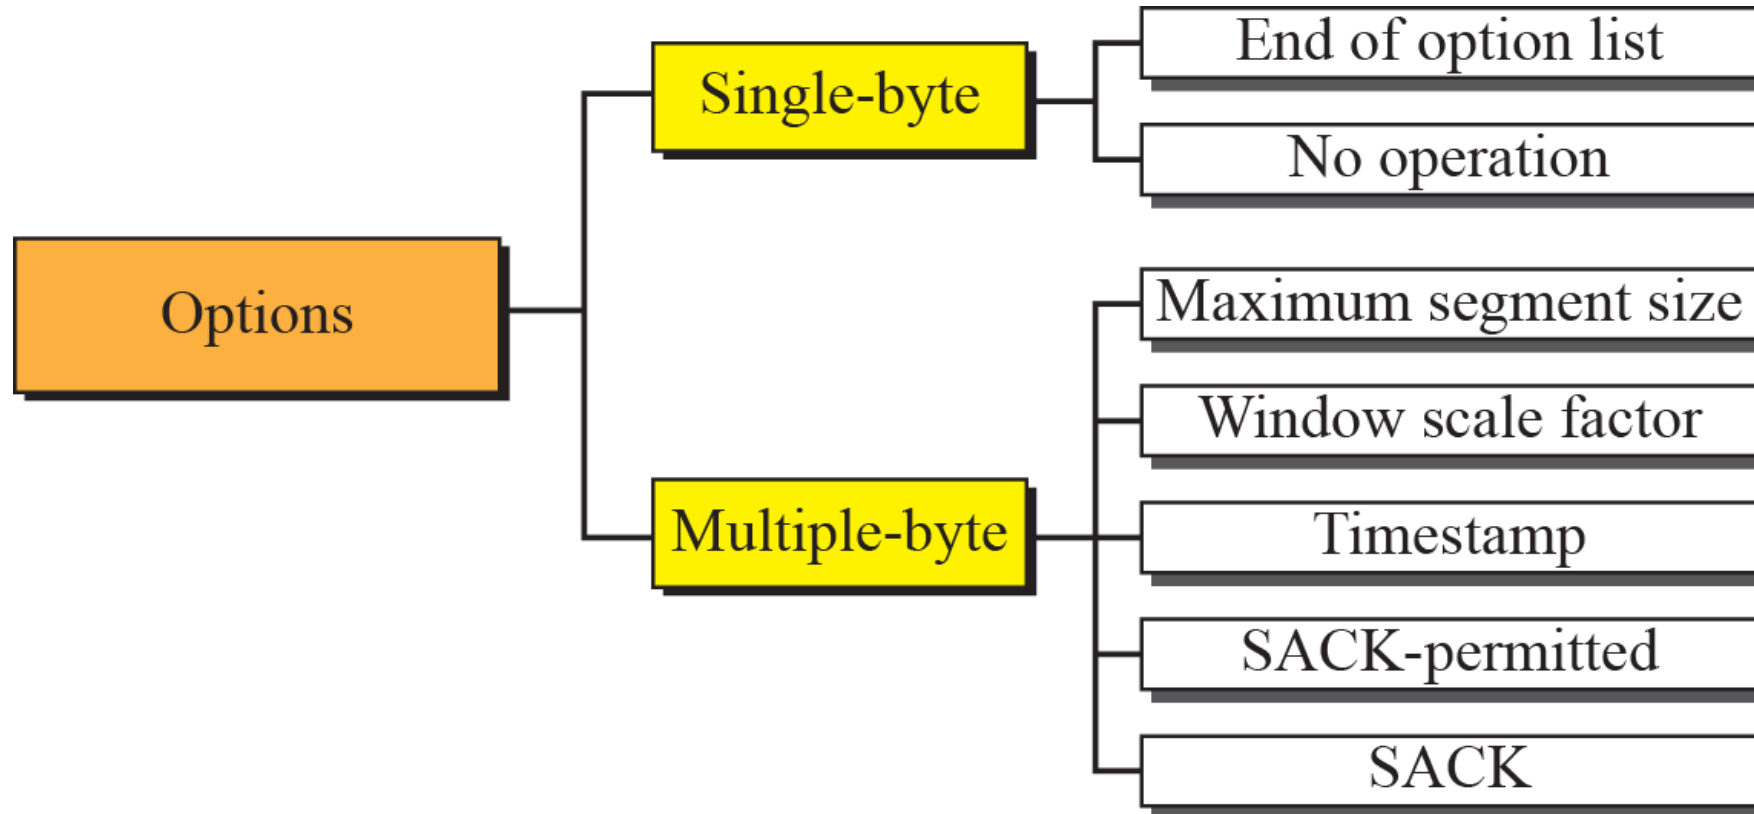


a. Segment

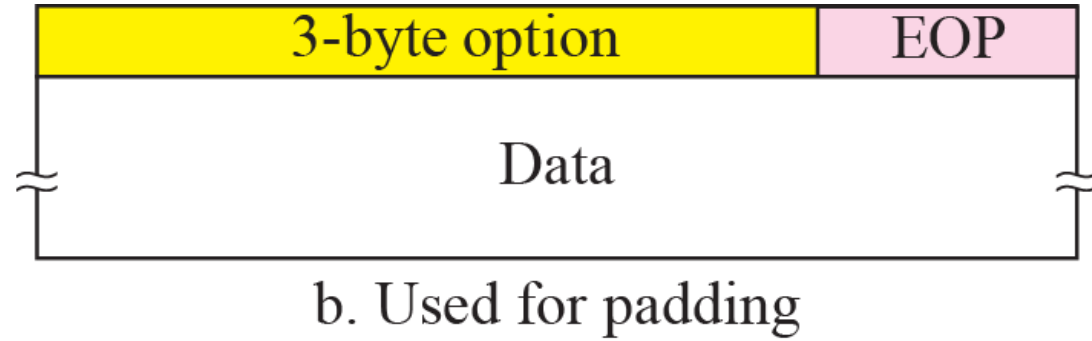



b. Header

## *Options*



## *End-of-option option*



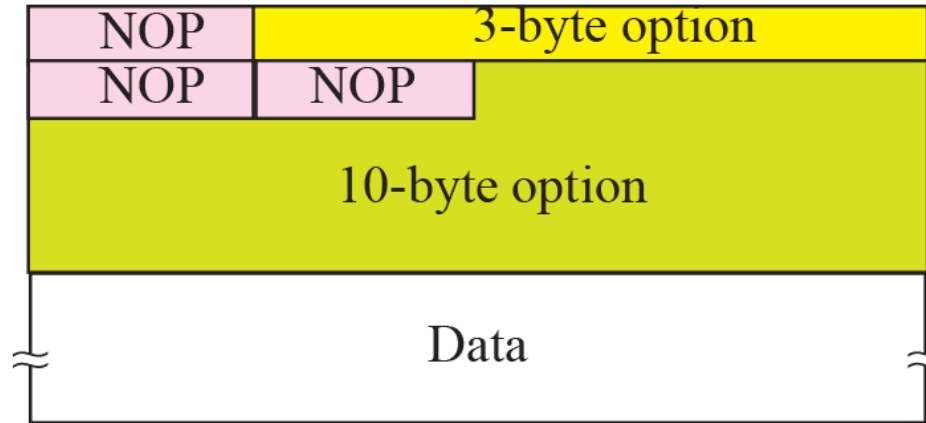


## *No-operation option*

- It is used to separate the different options used within the TCP Option field.



a. No operation option



b. Used to align beginning of an option

*NOP can be used more than once.*

## *Minimum-segment-size option*

- 1) The MSS (**Maximum Segment Size**) is defined as the **largest block of data** that a sender using TCP will send to the receiver.
- 2) When a connection is initiated a **SYN segment is sent**, in the process of sending a SYN segment, the sender has the **option of announcing its MSS**.
- 3) If a sender **doesn't** use the options field to declare the MSS then TCP assumes a **default of 536bytes** (minus the **20 byte TCP header**).


Kind: 2 00000010	Length: 4 00000100	Maximum segment size
1 byte	1 byte	2 bytes

*The value of MSS is determined during connection establishment and does not change during the connection.*

## Window-scale-factor option

1. For lines with high bandwidth, high delay or high bandwidth-delay product, 64KB is not enough to keep the sender busy all the time.
2. In RFC 1323, the use of the options field is proposed permitting the use of a window scale factor.
3. This scale factor permits shifting the window size up to 14 bits to the left therefore permitting window sizes of up to  $2^{(16+14)} = 2^{30}$  bytes.

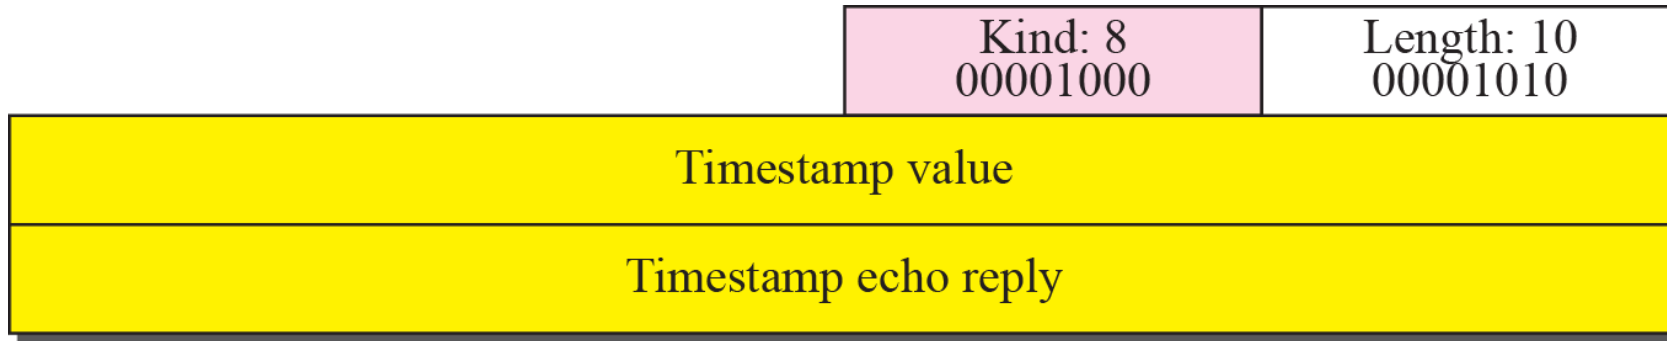
Example window size = 65535, and scale factor = 2 ( $2^2 = 4$  multiply by 4 to window size to calculate the actual window size)



*The value of the window scale factor can be determined only during connection establishment; it does not change during the connection.*

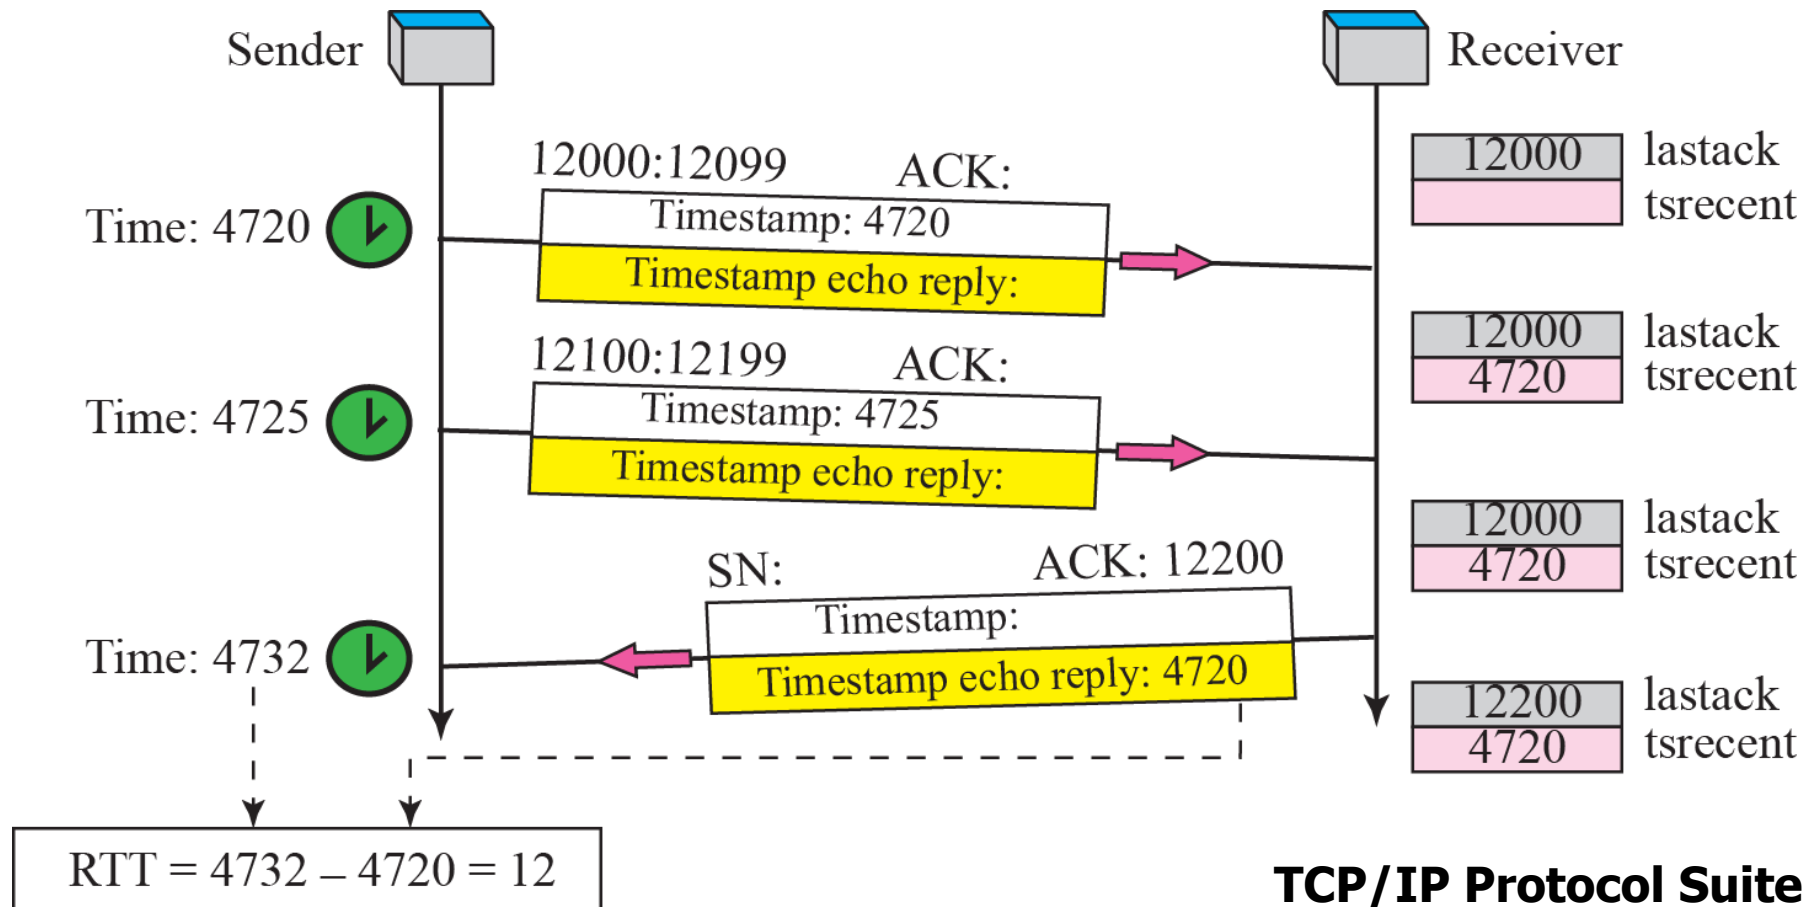
## *Timestamp option*

---

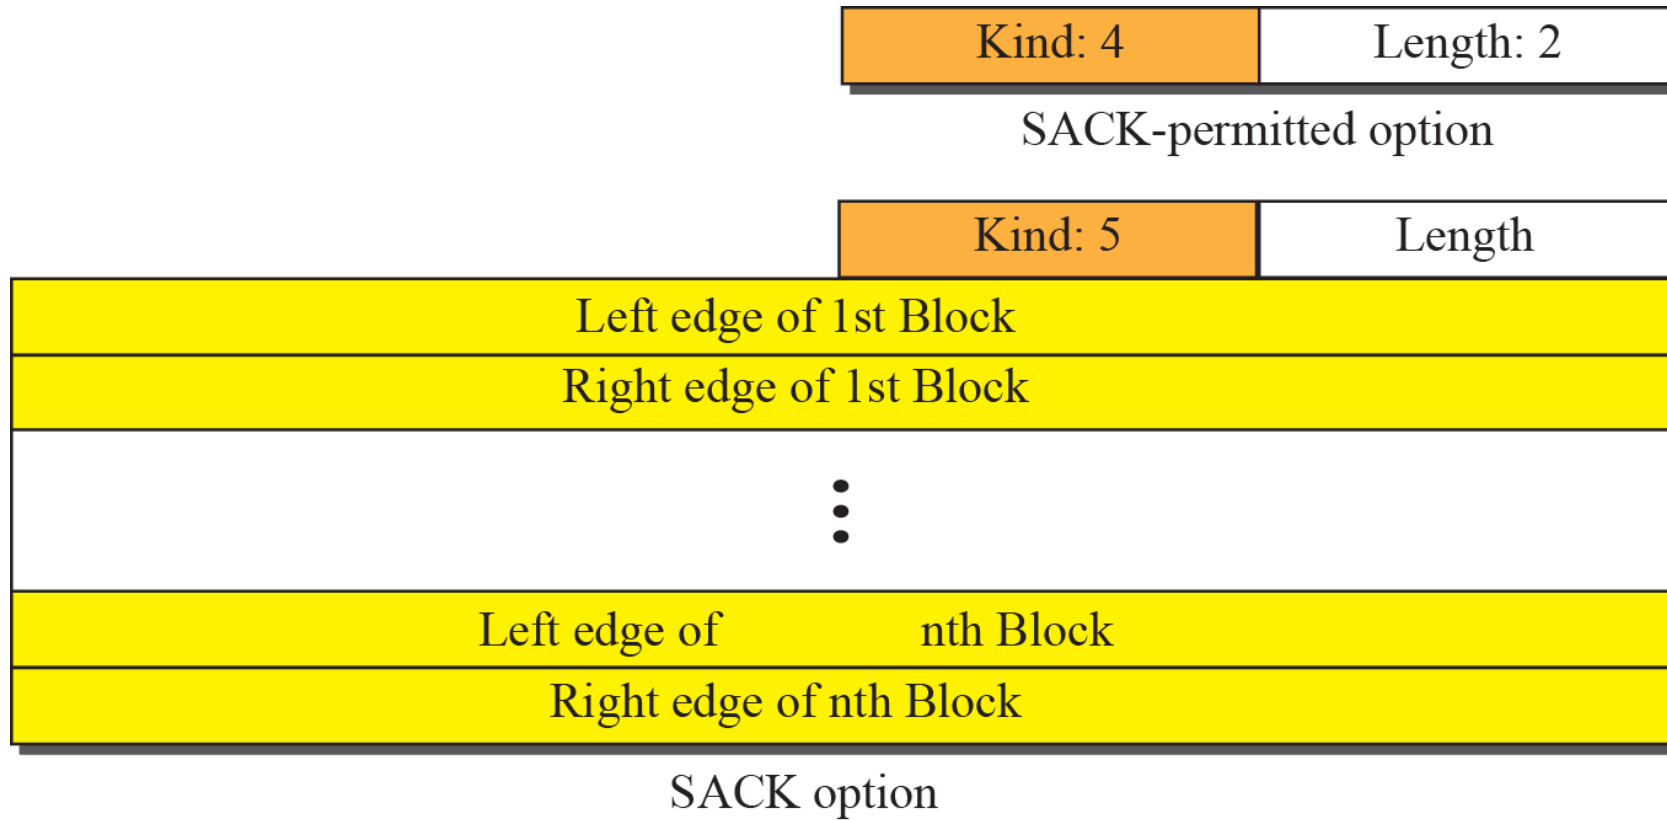


*One application of the timestamp option is the calculation of round-trip time (RTT).*

Figure shows an example that **calculates the round-trip time for one end**. Everything must be flipped if we want to calculate the RTT for the other end.



## Selective ACK (SACK)



Let us see how the **SACK option** is used to list out-of-order blocks. In Figure an end has received five segments of data.

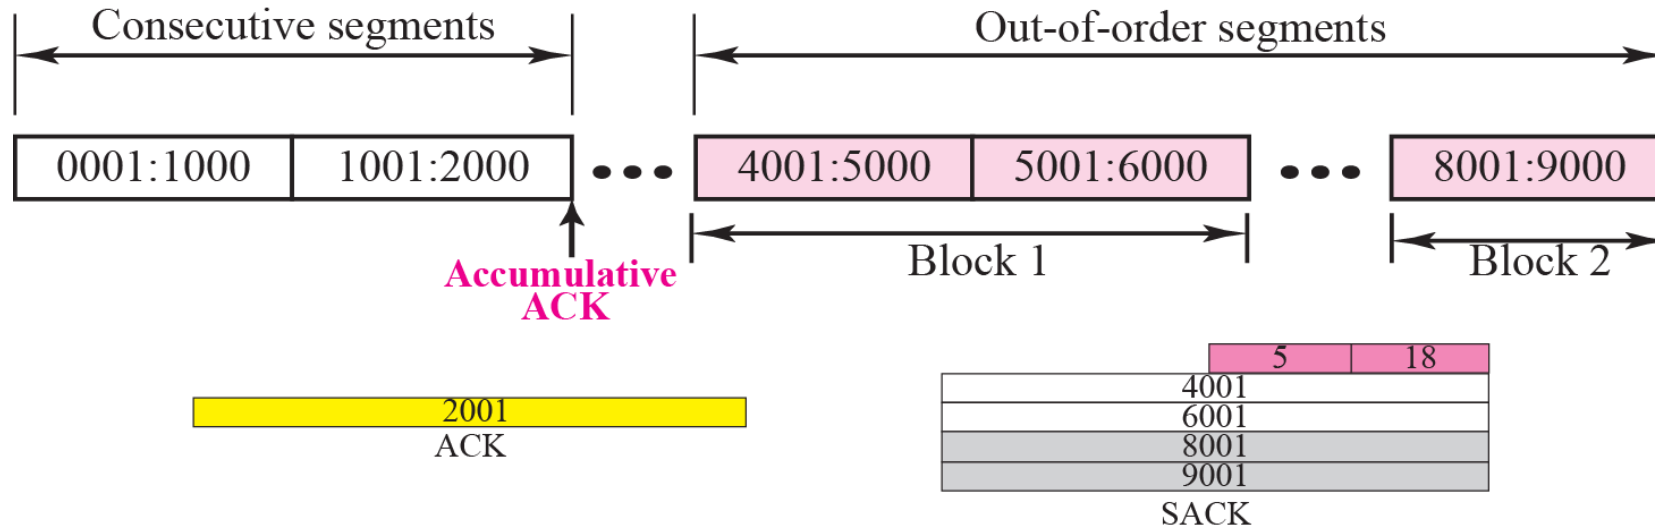


Figure shows how **a duplicate segment** can be detected with a combination of **ACK and SACK**.

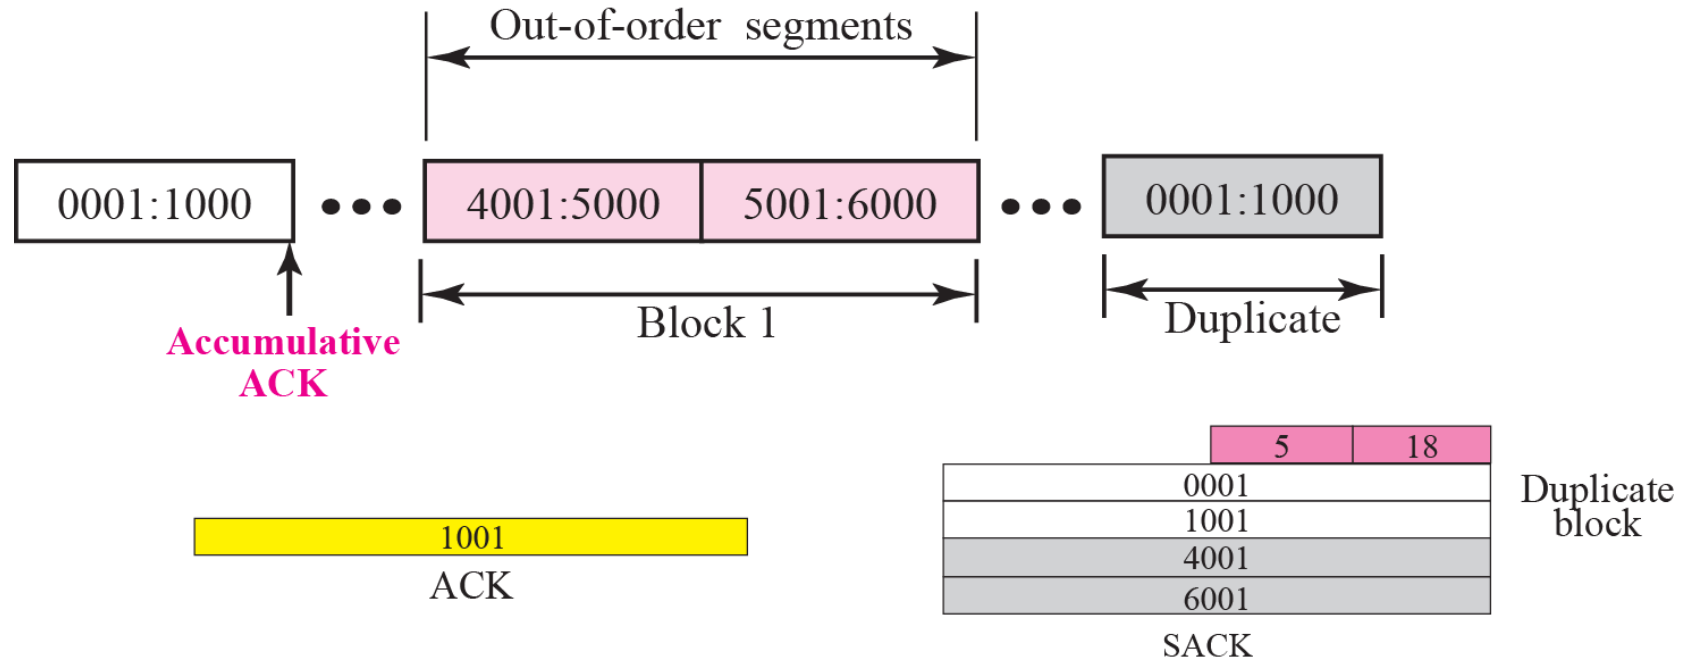
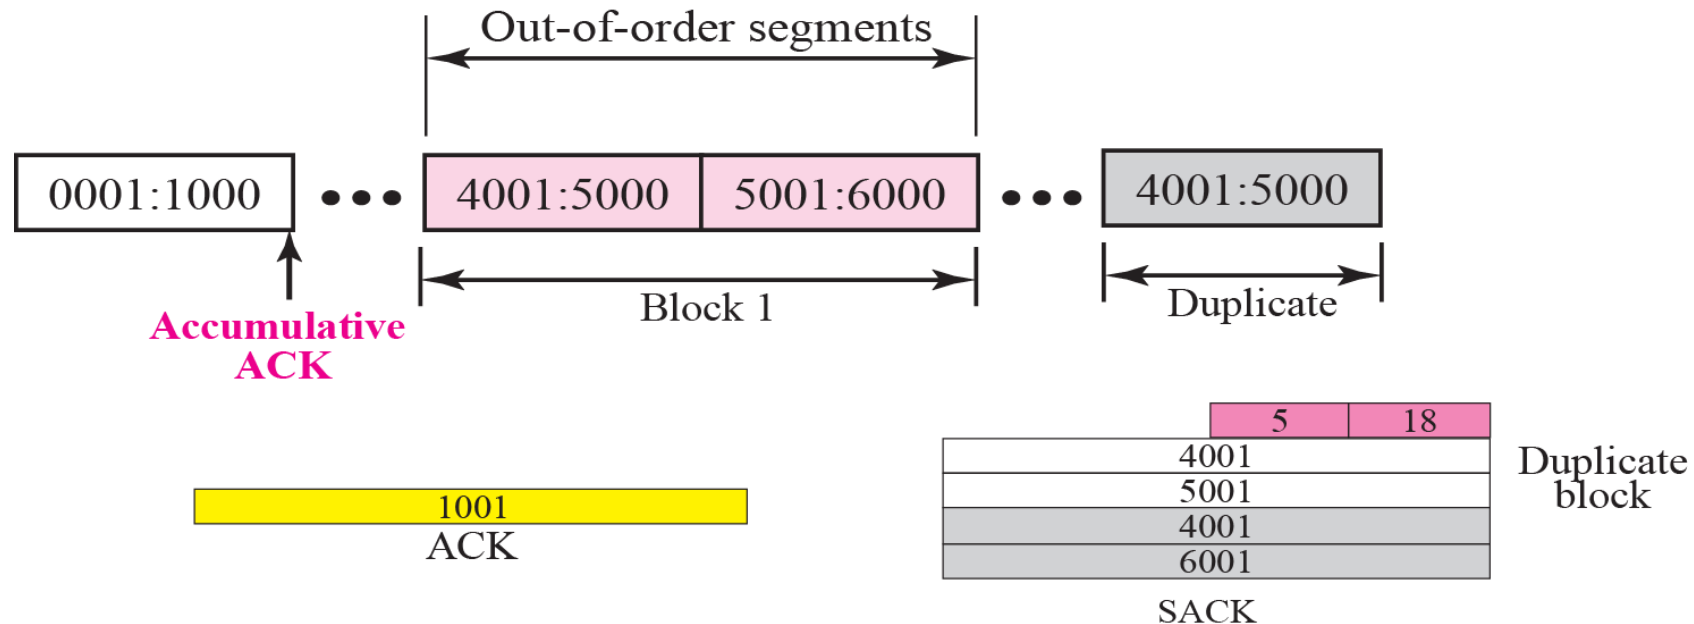




Figure shows what happens if one of the segments in the out-of-order section is also duplicated. In this example, one of the segments (4001:5000) is duplicated.

The SACK option announces this duplicate data first and then the out-of-order block.



# A New Transport Protocol



**QUIC: Quick UDP Internet Connections**

**HTTP/3: HTTP over QUIC is next Generation**


On 6 June 2022, IETF published HTTP/3 as a Proposed Standard in RFC 9114

# Introduction: Change

- Increasing scale of ..... everything
  - Flow size changes
  - Flow count increases (e.g., web pages)
  - Flow diversity increase (e.g., web pages)
    - Multiple connections



# HTTP Network Stack




TLS – Transport Layer Security  
TCP – Transport Control Protocol  
IP – Internet Protocol

## HTTP / 1.1

- January 1997
- **Many parallel TCP connection (6 connections per host name)**
- **HTTP head of line blocking**


## HTTP / 2

- May 2015
- **Using Single connection per host**
- **Many parallel streams**
- **TCP head of line blocking**

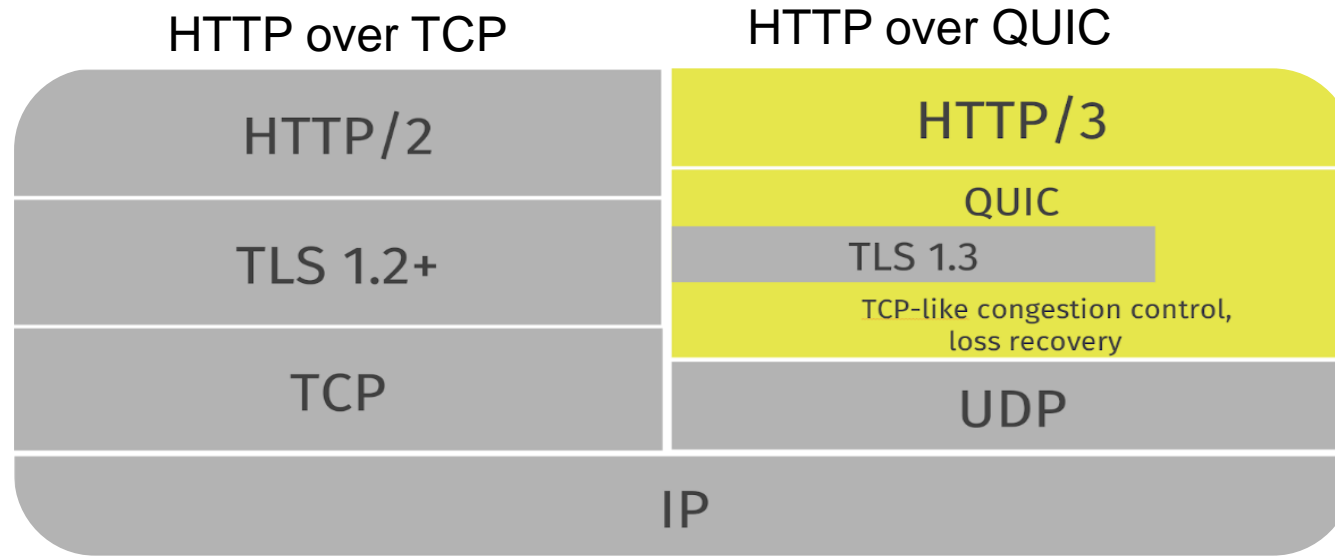


# QUIC: Quick UDP Internet Connections

- application-layer protocol, on top of UDP
  - increase performance of HTTP
  - deployed on many Google servers, apps (Chrome, mobile YouTube app)




# HTTP Over QUIC Network Stack



HTTP / 3

Workgroup: QUIC  
Internet-Draft: draft-ietf-quic-http-latest  
Published: 9 June 2020  
Intended Status: Standards Track  
Expires: 11 December 2020  
Author: M. Bishop, Ed.  
Akamai

- **No - TCP head of line blocking**
  - **streams are independent to each other**
- **Faster handshake**
  - **Earlier data**
- **More encryption, always**
- **Over UDP** (*Connection less, No resend, No flow control*)




- Stream impact not connection impact

# QUIC: Quick UDP Internet Connections

adopts approaches we've studied in this chapter for connection establishment, error control, congestion control


- **error and congestion control:** “Readers familiar with TCP’s loss detection and congestion control will find algorithms here that parallel well-known TCP ones.” [from QUIC specification]
- **connection establishment:** reliability, congestion control, authentication, encryption, state established in one RTT
- multiple application-level “streams” multiplexed over single QUIC connection
  - separate reliable data transfer, security
  - common congestion control

# QUIC: Connection establishment



TCP (reliability, congestion control state) + TLS (authentication, crypto state)

- 2 serial handshakes





QUIC: reliability, congestion control, authentication, crypto state

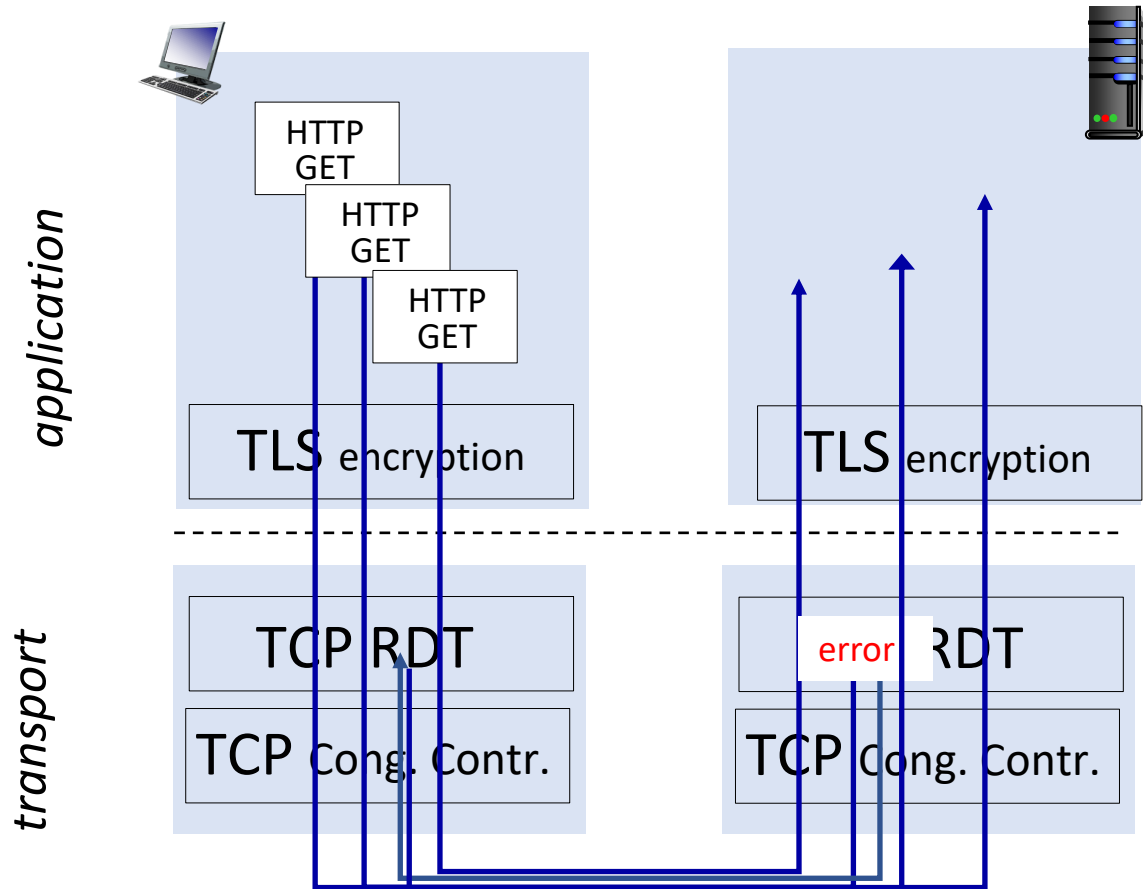
- 1 handshake



# Subsequent Connection to the same server



# QUIC: streams: parallelism, no HOL blocking



(a) HTTP 1.1

# QUIC Status

- In May 2021, the IETF standardized QUIC in [RFC 9000](#).

## Implementations:

Apple, Facebook, Fastly, Firefox, F5, Google, Microsoft ...

## Server deployments have been going on for a while

Akamai, Cloudflare, Facebook, Fastly, Google ...

## Clients are at different stages of deployment

Chrome, Firefox, Edge, Safari  
iOS, MacOS